

APPLICATIONS



OF DATA SCIENCE

The Pandasverse

Applications of Data Science - Class 6

Giora Simchoni

`gsimchoni@gmail.com` and add `#dsapps` in subject

Stat. and OR Department, TAU

2020-03-01

APPLICATIONS



OF DATA SCIENCE

Numpy: Your best friend

APPLICATIONS



OF DATA SCIENCE

Python was not made for Data Science

```
mean([1, 2, 3, 4, 5])
```

```
## Error in py_call_impl(callable, dots$args, dots$keywords): NameError: name  
'mean' is not defined  
##  
## Detailed traceback:  
## File "", line 1, in
```

Enter Numpy to the rescue:

```
import numpy as np  
  
np.mean(np.array([1, 2, 3, 4, 5]))
```

```
## 3.0
```

```
np.array([1, 2, 3, 4, 5]).mean()
```

```
## 3.0
```

Numpy Arrays

Create with a list:

```
a = np.array([1, 2, 3])  
print(type(a))
```

```
## <class 'numpy.ndarray'>
```

```
print(a.shape)
```

```
## (3,)
```

⚠ Index is zero based!

```
print(a[0])
```

```
## 1
```

Create a 2D array:

```
b = np.array([[1,2,3],[4,5,6]])  
print(b)
```

```
## [[1 2 3]  
##  [4 5 6]]
```

```
print(b.shape)
```

```
## (2, 3)
```

Many ways to create "typical" arrays:

```
# create an array of all zeros  
# (the parameter is a tuple specifying the array shape)  
a = np.zeros((2,2))  
  
# create an array of all ones  
b = np.ones((1,2))  
  
# create a constant array  
c = np.full((2,2), 7)  
  
# create a 2x2 identity matrix  
d = np.eye(2)  
  
# create an array filled with random U(0, 1) values  
e = np.random.random((2,2))  
  
# create a sequence from 2 to 15, not including  
np.arange(2, 15)  
  
# create sequence of 11 numbers between 0 and 1 including  
np.linspace(0, 1, 11)
```

And every array has a `reshape()` method:

```
np.arange(0.1, 1, step=0.1).reshape(3, 3)
```

```
## array([[0.1, 0.2, 0.3],  
##        [0.4, 0.5, 0.6],  
##        [0.7, 0.8, 0.9]])
```


Numpy Math

Elementwise multiplication:

```
x = np.array([[1,2],[3,4]], dtype=np.float64)
y = np.array([[5,6],[7,8]], dtype=np.float64)
print(x * 2)
```

```
## [[2. 4.]
##  [6. 8.]]
```

Elementwise sum:

```
print(x + y)
```

```
## [[ 6.  8.]
##  [10. 12.]]
```

Same:

```
print(np.add(x, y))
```

You get the idea:

```
print(x - y)
print(np.subtract(x, y))

print(x * y)
print(np.multiply(x, y))

print(x / y)
print(np.divide(x, y))

print(np.sqrt(x))
```

Vector/Matrix multiplication:

```
print(x.dot(y))
```

```
## [[19. 22.]  
##  [43. 50.]]
```

```
print(np.dot(x, y))
```

```
## [[19. 22.]  
##  [43. 50.]]
```

```
v = np.array([9, 10])  
w = np.array([11, 12])  
  
print(v.dot(w))
```

```
## 219
```

```
print(np.dot(v, w))
```

```
## 219
```

Transpose

```
x = np.array([[1,2],[3,4]])  
print(x.T)
```

```
## [[1 3]  
##  [2 4]]
```

Sum, mean, std, median, quantile, min, max...:

```
print(np.sum(x))  # Compute sum of all elements
```

```
## 10
```

```
print(np.sum(x, axis=0))  # Compute sum of each column
```

```
## [4 6]
```

```
print(np.std([1,2,3]))  # possible, in case you were wondering
```

```
## 0.816496580927726
```

Numpy Indexing and Slicing

Similar to R but there are some things worth noticing:

```
a = np.array([[1,2,3,4], [5,6,7,8], [9,10,11,12]])  
print(a)
```

```
## [[ 1  2  3  4]  
##   [ 5  6  7  8]  
##   [ 9 10 11 12]]
```

```
# use slicing to pull out the subarray consisting of the first 2 rows  
# and columns 1 and 2; b of shape (2, 2)  
b = a[:2, 1:3]  
print(b)
```

```
## [[2 3]  
##   [6 7]]
```

```
# a slice of an array is a view into the same data, so modifying :  
# will modify the original array.  
print(a[0, 1])
```

```
## 2
```

```
b[0, 0] = 77  
print(a[0, 1])
```

```
## 77
```

Very convenient, R does not have these features without external packages:

```
# index "from last place"  
a[-2:]
```

```
## array([[ 5,  6,  7,  8],  
##       [ 9, 10, 11, 12]])
```

```
# reverse an array  
a = np.arange(5)  
print(a[::-1])
```

```
## [4 3 2 1 0]
```

Working with boolean masks like in R:

```
print(a[a > 2])
```

```
## [3 4]
```

```
print(a[np.where(a > 2)])
```

```
## [3 4]
```

```
print(a[np.argmin(a)])
```

```
## 0
```


Scipy: Scientific Computing and Stats

APPLICATIONS



OF DATA SCIENCE

Many modules, let's focus on:

- **sparse**: Sparse Matrices manipulation
- **ndimage**: Images manipulation (though see `scikit-image` and `opencv`)
- **stats**: Statistics (though see `statsmodels`)

sparse

```
from scipy.sparse import csr_matrix

row = np.array([0, 0, 1, 2, 2, 2])
col = np.array([0, 2, 2, 0, 1, 2])
data = np.array([1, 2, 3, 4, 5, 6])
sparse_a = csr_matrix((data, (row, col)), shape=(3, 3))

print(sparse_a.toarray())
```

```
## [[1 0 2]
##  [0 0 3]
##  [4 5 6]]
```

ndimage

```
from scipy import ndimage
from scipy import misc
import matplotlib.pyplot as plt

face = misc.face(gray=True)
blurred_face = ndimage.gaussian_filter(face, sigma=10)


print(face.shape)
```

```
## (768, 1024)
```

```
print(face[:5, :5])
```

```
## [[114 130 145 147 147]
##   [ 83 104 123 130 134]
##   [ 68  88 109 116 120]
##   [ 78  94 109 116 121]
##   [ 99 109 119 128 138]]
```

```
plt.subplot(121)
plt.imshow(face, cmap=plt.cm.gray)
plt.subplot(122)
plt.imshow(blurred_face, cmap=plt.cm.gray)
plt.show()
```



stats

```
from scipy import stats
```

```
rvs1 = stats.norm.rvs(loc=5,scale=10,size=500)
```

```
rvs2 = stats.norm.rvs(loc=5,scale=10,size=500)
```

```
stats.ttest_ind(rvs1,rvs2)
```

```
## Ttest_indResult(statistic=-0.1672020280648639, pvalue=0.8672449606209668)
```

Pandas: Data, Data, Data

APPLICATIONS



OF DATA SCIENCE

Series, DataFrames

A Pandas Series is a vector of data, a column.

```
import pandas as pd

s = pd.Series([1,3,5,np.nan,6,8])
print(s)
```

```
## 0    1.0
## 1    3.0
## 2    5.0
## 3    NaN
## 4    6.0
## 5    8.0
## dtype: float64
```


A DataFrame is a data table, always indexed.

Creating one from a random numpy 2D array (notice the index isn't specified, automatically becomes zero based counter):

```
df = pd.DataFrame(np.random.randn(6,4), columns = ['A', 'B', 'C',  
print(df)
```

##		A	B	C	D
## 0		0.637024	0.328160	-0.893112	0.924013
## 1		1.659861	-0.531544	0.125794	-1.574434
## 2		0.029313	-0.394294	-0.708947	-0.046269
## 3		-0.719662	0.374894	2.028673	-0.751812
## 4		-0.969941	0.287350	-0.462109	0.971948
## 5		0.306526	-1.873495	0.965911	-1.425832

Creating a DataFrame from a very varied dictionary where each key is a column (also see `pd.from_dict()`).

```
df2 = pd.DataFrame({'A' : 1.,
                    'B' : pd.Timestamp('20130102'),
                    'C' : pd.Series(1, index = list(range(4)),
                                   dtype = 'float32'),
                    'D' : np.array(np.arange(4), dtype = 'int32'),
                    'E' : pd.Categorical(
                        ["test", "train", "test", "train"]
                    ),
                    'F' : 'foo' })

print(df2)
```

##	A	B	C	D	E	F
## 0	1.0	2013-01-02	1.0	0	test	foo
## 1	1.0	2013-01-02	1.0	1	train	foo
## 2	1.0	2013-01-02	1.0	2	test	foo
## 3	1.0	2013-01-02	1.0	3	train	foo

read_csv()

```
okcupid = pd.read_csv("../data/okcupid.csv.zip")
```

```
okcupid.shape
```

```
## (59946, 31)
```

```
okcupid.columns
```

```
## Index(['age', 'body_type', 'diet', 'drinks', 'drugs', 'education', 'essa  
##      'essay1', 'essay2', 'essay3', 'essay4', 'essay5', 'essay6', 'essa  
##      'essay8', 'essay9', 'ethnicity', 'height', 'income', 'job',  
##      'last_online', 'location', 'offspring', 'orientation', 'pets',  
##      'religion', 'sex', 'sign', 'smokes', 'speaks', 'status'],  
##      dtype='object')
```

info(), describe(), head() and tail()

```
okcupid.describe()
```

```
##              age              height              income
## count    59946.000000    59943.000000    59946.000000
## mean       32.340290       68.295281     20033.222534
## std        9.452779        3.994803     97346.192104
## min        18.000000        1.000000       -1.000000
## 25%        26.000000       66.000000       -1.000000
## 50%        30.000000       68.000000       -1.000000
## 75%        37.000000       71.000000       -1.000000
## max       110.000000       95.000000    1000000.000000
```

```
okcupid.head(3)
```

```
##    age  ...    status
## 0   22  ...    single
## 1   35  ...    single
## 2   38  ...  available
##
## [3 rows x 31 columns]
```

Not `data.frame`, `DataFrame`

dplyr	pandas
<code>mutate</code>	<code>assign</code>
<code>select</code>	<code>filter</code>
<code>rename</code>	<code>rename</code>
<code>filter</code>	<code>query</code>
<code>arrange</code>	<code>sort_values</code>
<code>group_by</code>	<code>groupby</code>
<code>summarize</code>	<code>agg</code>

💡 There *are* Pandas dialects, don't go translating your pipes verbatim.

assign()

Add a column `height_cm`, the height in centimeters:

```
okcupid = okcupid.assign(height_cm = okcupid['height'] * 2.54)
okcupid = okcupid.assign(height_cm = lambda x: x.height * 2.54)
```

If you don't need a pipe just do:

```
okcupid['height_cm'] = okcupid['height'] * 2.54
```

query() and filter()

Query only women, filter only age and height:

```
okcupid \  
  .query('sex == "f") \  
  .filter(['age', 'height']) \  
  .head(5)
```

	age	height
## 6	32	65.0
## 7	31	65.0
## 8	24	67.0
## 13	30	66.0
## 14	29	62.0

Again, without a pipe:

```
okcupid[okcupid['sex'] == "f"][['age', 'height']]
```

Same but income over 100K, and select all essay questions:

```
okcupid \  
  .query('sex == "f" and income > 100000') \  
  .filter(okcupid.columns[okcupid.columns.str.startswith('essay')])
```

```
##                                     essay0 ...  
## 48      i love it here, except when it's hotter than a... ... if you da  
## 188     i'm silly. i'm analytical. i'm fond of short s... ... you want  
## 301     welcome... i am one genuine, straight forward,... ...  
## 337     purebred cali girl! born and raised in nor cal... ... you are a  
## 402     i wasn't like every other kid, you know, who d... ... you think  
## ...                                     ... ...  
## 59326   i am a forensic psychologist, mother, sister a... ...  
## 59395                                     NaN ...  
## 59789   i'm a fun loving woman, romantic, faithful, ea... ...  
## 59818   hello, i am usually pretty shy and sometimes a... ... you are p  
## 59819   this is a pretty good read. admittedly windy. ... ... you like  
##  
## [208 rows x 10 columns]
```


agg()

Find the average height of women

```
okcupid \
    .query('sex == "f"') \
    .filter(['height_cm']) \
    .agg('mean')
```

```
## height_cm      165.363837
## dtype: float64
```

Notice we got a `pd.Series`, the Pandas equivalent for a vector. We could use the `.values` attribute to pull the Numpy array behind the Series:

```
okcupid \
    .query('sex == "f"') \
    .filter(['height_cm']) \
    .agg('mean').values
```

```
## array([165.36383729])
```

groupby()

But why settle for women only?

```
okcupid \
.groupby('sex')['height_cm'] \
.agg('mean')
```

```
## sex
## f      165.363837
## m      178.926471
## Name: height_cm, dtype: float64
```

And you might want to consider `rename()` ing sex!

```
okcupid \
.groupby('sex')['height_cm'] \
.agg('mean') \
.rename_axis(index = {'sex': 'gender'})
```

```
## gender
## f      165.363837
## m      178.926471
## Name: height_cm, dtype: float64
```

Group by multiple variables, get more summaries, arrange by descending average height:

```
okcupid \
.groupby(['sex', 'status'])['height_cm'] \
.agg(['mean', 'median', 'count']) \
.sort_values('median', ascending=False)
```

		mean	median	count
##	sex status			
##	m available	179.445012	180.34	1209
##	married	179.454629	180.34	175
##	seeing someone	179.257926	177.80	1061
##	single	178.894660	177.80	33376
##	unknown	177.376667	176.53	6
##	f available	166.381616	166.37	656
##	married	165.871407	165.10	135
##	seeing someone	165.431745	165.10	1003
##	single	165.328643	165.10	22318
##	unknown	160.655000	158.75	4

Pro tip: size ()

When all you want is, well, size:

```
okcupid.groupby('body_type').size()
```

```
## body_type
## a little extra      2629
## athletic           11819
## average             14652
## curvy               3924
## fit                 12711
## full figured        1009
## jacked              421
## overweight          444
## rather not say      198
## skinny              1777
## thin                4711
## used up             355
## dtype: int64
```

loc, iloc and at

loc is for selection by name:

```
okcupid.loc[:3, ['sex', 'height_cm']]
```

```
##      sex  height_cm
## 0      m    190.50
## 1      m    177.80
## 2      m    172.72
## 3      m    180.34
```

The first element to loc slices the index by name. The reason that ":3" works is that our index is numeric. If it were for example ['a', 'b', 'c', ...] it would not have worked.

loc can also accept boolean indexing:

```
okcupid.loc[okcupid['sex'] == 'm', 'height_cm']
```

`iloc` is for selection by integers on the index or column indices

```
okcupid.iloc[:3, 1:3]
```

```
##          body_type          diet
## 0  a little extra  strictly anything
## 1          average      mostly other
## 2             thin      anything
```

This would have worked also if the index was `['a', 'b', 'c', ...]`.

Finally `at` is for accessing a specific value fast:

```
okcupid.at[1989, 'body_type']
```

```
## 'average'
```

seaborn: Visualization


APPLICATIONS



OF DATA SCIENCE

```
import matplotlib.pyplot as plt
import seaborn as sns

sns.set()
g = sns.scatterplot('age', 'height_cm', hue='sex', data = okcupid)
plt.show()
```




```
g = sns.relplot('age', 'height_cm',  
                hue = 'sex', kind = 'scatter', col='sex', data = okcupid)  
plt.show()
```

