

APPLICATIONS



OF DATA SCIENCE

The Trees

Applications of Data Science - Class 11

Giora Simchoni

gsimchoni@gmail.com and add #dsapps in subject

Stat. and OR Department, TAU

2020-02-09

APPLICATIONS



OF DATA SCIENCE

The Pros and Cons of Trees

APPLICATIONS



OF DATA SCIENCE

Pros

- It's just a set of if-else statements my 7 y/o can get
- Highly interpretable (when tree is not large)
- Easy to implement
- Fast? (to predict)
- Little pre-processing of predictors needed
- Handle all types of predictors (continuous, categorical)
- Handle missing data, *predict* missing data
- Assumption free
- Feature selection built-in (but when predictors are correlated...)
- Low bias, in general

Cons

- High variance, in general
- Rectangular predictor regions - not always a good thing
- Complexity of prediction limited in no. of leaves! (For a simple CART)
- Not so fast? (to train)
- Greedy
- Selection Bias of predictors with more distinct values?
- Variable Importance when predictors are correlated

Detour: A Regression Problem

APPLICATIONS



OF DATA SCIENCE

OKCupid: Predicting Annual Income

It won't be easy:

```
okcupid <- read_csv("../data/okcupid.csv.zip")  
okcupid %>% count(income, sort = TRUE) %>% head(20)
```

```
## # A tibble: 13 x 2  
##       income     n  
##       <dbl> <int>  
## 1      -1  48442  
## 2    20000   2952  
## 3   100000   1621  
## 4    80000   1111  
## 5    30000   1048  
## 6    40000   1005  
## 7    50000    975  
## 8    60000    736  
## 9    70000    707  
## 10   150000    631  
## 11  1000000    521  
## 12   250000    149  
## 13   500000     48
```

We will stick to non-NA (income) observations, and predict $\log_{10}(income/100000)$:

```
okupid2 <- okupid %>%
  mutate(income = ifelse(income == -1, NA, log10(income/100000)))
  drop_na(income)
```

In the vector `predictors` (see slides Rmd source) we have 42 continuous and categorical variables which may or may not be predictive to income:

```
okupid2 <- okupid2 %>%
  select(income, predictors) %>%
  mutate(id = 1:n())

dim(okupid2)
```

```
## [1] 11504     44
```

```
glimpse(okcupid2)
```

```
## Observations: 11,504
## Variables: 44
## $ income <dbl> -0.09691001, -0.69897000, -0.39794001, ...
## $ age <dbl> 35, 23, 28, 30, 29, 40, 31, 22, 35, 31, ...
## $ height_cm <dbl> 177.80, 180.34, 182.88, 167.64, 157.48, ...
## $ sex <fct> m, m, m, f, f, m, f, m, m, f, m, m, f, m...
## $ body_type <fct> average, thin, average, skinny, thin, fi...
## $ body_type_not_perfect <fct> TRUE, FALSE, TRUE, FALSE, FALSE, FALSE, ...
## $ diet2 <fct> other, vegetarian, anything, anything, a...
## $ drinks <fct> often, socially, socially, socially, soc...
## $ drugs <fct> sometimes, NA, never, never, never, NA, ...
## $ religion2 <fct> atheist, NA, christian, christian, chris...
## $ education2 <fct> other, student1, degree1, high_school, s...
## $ education_kind <fct> working, working, graduated, graduated, ...
## $ education_academic <fct> FALSE, FALSE, TRUE, FALSE, FALSE, TRUE, ...
## $ ethnicity2 <fct> white, white, white, white, other, white...
## $ part_black <fct> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE...
## $ part_white <fct> TRUE, TRUE, TRUE, TRUE, TRUE, NA, ...
## $ part_asian <fct> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE...
## $ part_hispanic <fct> FALSE, FALSE, FALSE, FALSE, TRUE, FALSE, ...
## $ job3 <fct> travel, student, financial, marketing, o...
## $ orientation <fct> straight, straight, straight, straight, ...
## $ pets_has_dogs <fct> FALSE, FALSE, FALSE, TRUE, FALSE, FALSE, ...
## $ pets_has_cats <fct> FALSE, FALSE, FALSE, FALSE, TRUE, FALSE, ...
## $ pets_likes_cats <fct> TRUE, TRUE, TRUE, TRUE, FALSE, FALSE, NA...
## $ pets_likes_dogs <fct> TRUE, FALSE, FALSE, FALSE, TRUE, TRUE, N...
## $ sian_fun <fct> FALSE, FALSE, FALSE, NA, FALSE, TRUE, NA...
```

We split the data into training, validation and test sets:

```
# test_idx <- sample(1:nrow(okcupid2), 2000, replace = FALSE)
# train_idx <- okcupid2 %>% filter(!id %in% test_idx) %>% sample_
# valid_idx <- okcupid2 %>% filter(!id %in% test_idx, !id %in% tra
okcupid2 <- okcupid2 %>% select(-id)

idx <- read_rds("../data/okcupid2_idx.rda")
train_idx <- idx$train_idx
valid_idx <- idx$valid_idx
test_idx <- idx$test_idx

okcupid2_train <- okcupid2[train_idx, ]
okcupid2_valid <- okcupid2[valid_idx, ]
okcupid2_test <- okcupid2[test_idx, ]


library(glue)
glue("train no. of rows: {nrow(okcupid2_train)}
      validation no. of rows: {nrow(okcupid2_valid)}
      test no. of rows: {nrow(okcupid2_test)}")
```



```
## train no. of rows: 8000
## validation no. of rows: 1504
## test no. of rows: 2000
```

Our transformed income dependent variable behaves "ok":


```
ggplot(okcupid2_train, aes(income)) +  
  geom_histogram(fill = "red", alpha = 0.5) +  
  theme_classic()
```



We can quickly see percentage of missing values with [naniar](#):

```
library(naniar)

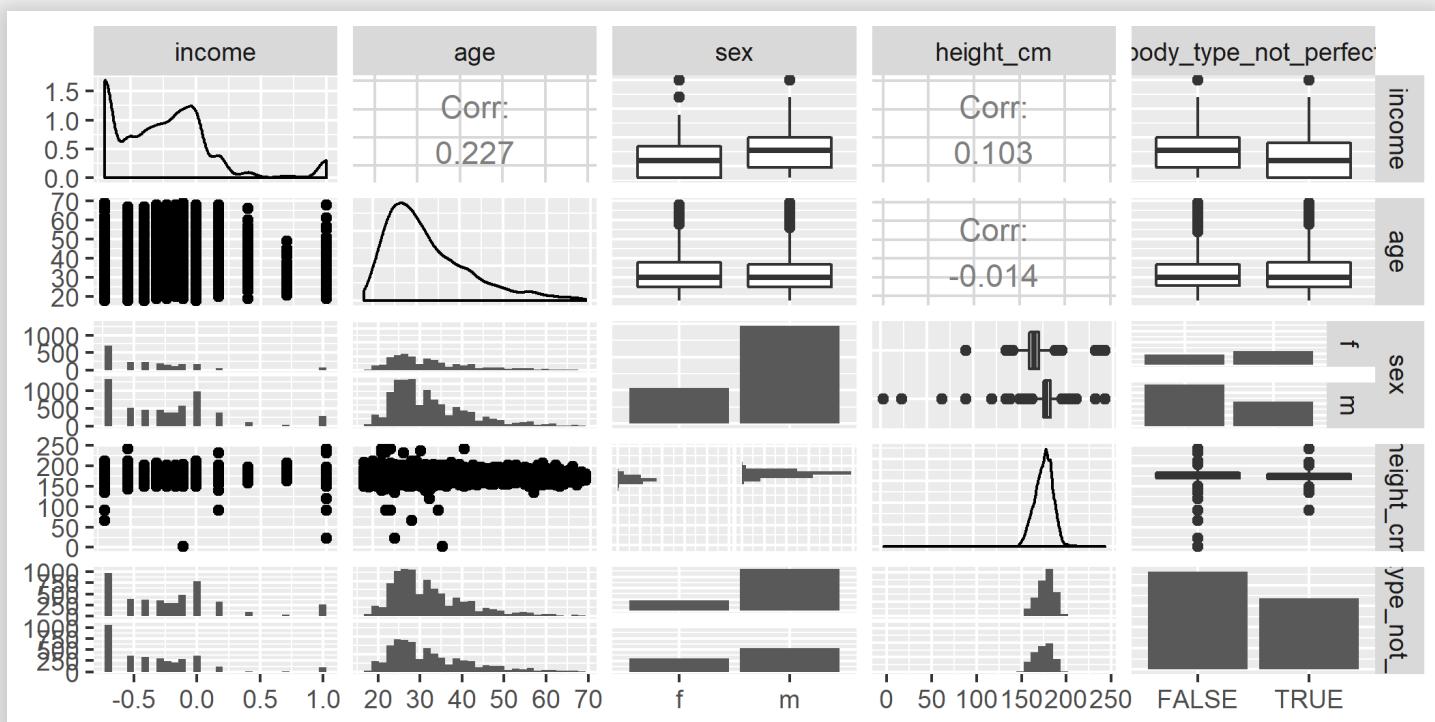
vis_miss(okcupid2_train %>%
          sample_frac(0.2) %>%
          select(-starts_with("essay")))
```



Also worth exploring some basic relations between predictors and income. You can use the work of others, e.g. [ggpairs](#):

```
library(GGally)
```


```
ggpairs(okcupid2_train %>%
           select(income, age, sex, height_cm, body_type_not_perfec
```



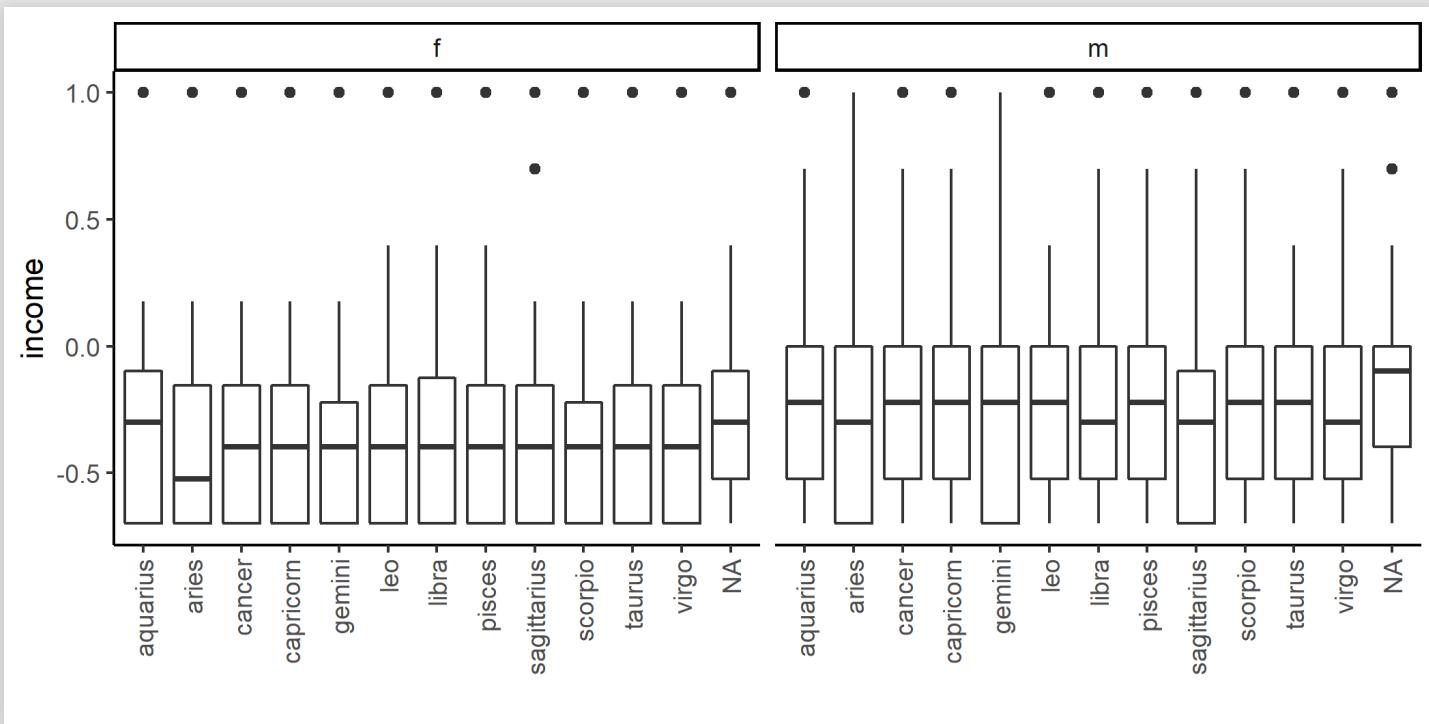
But don't be ashamed of simply exploring on your own:

```
var_vs_income_boxplot <- function(var) {  
  ggplot(okcupid2_train, aes({{var}}, income)) +  
  geom_boxplot() +  
  facet_wrap(~ sex) +  
  theme_classic() +  
  labs(x = "") +  
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust =  
})
```


```
var_vs_income_boxplot(body_type)
```



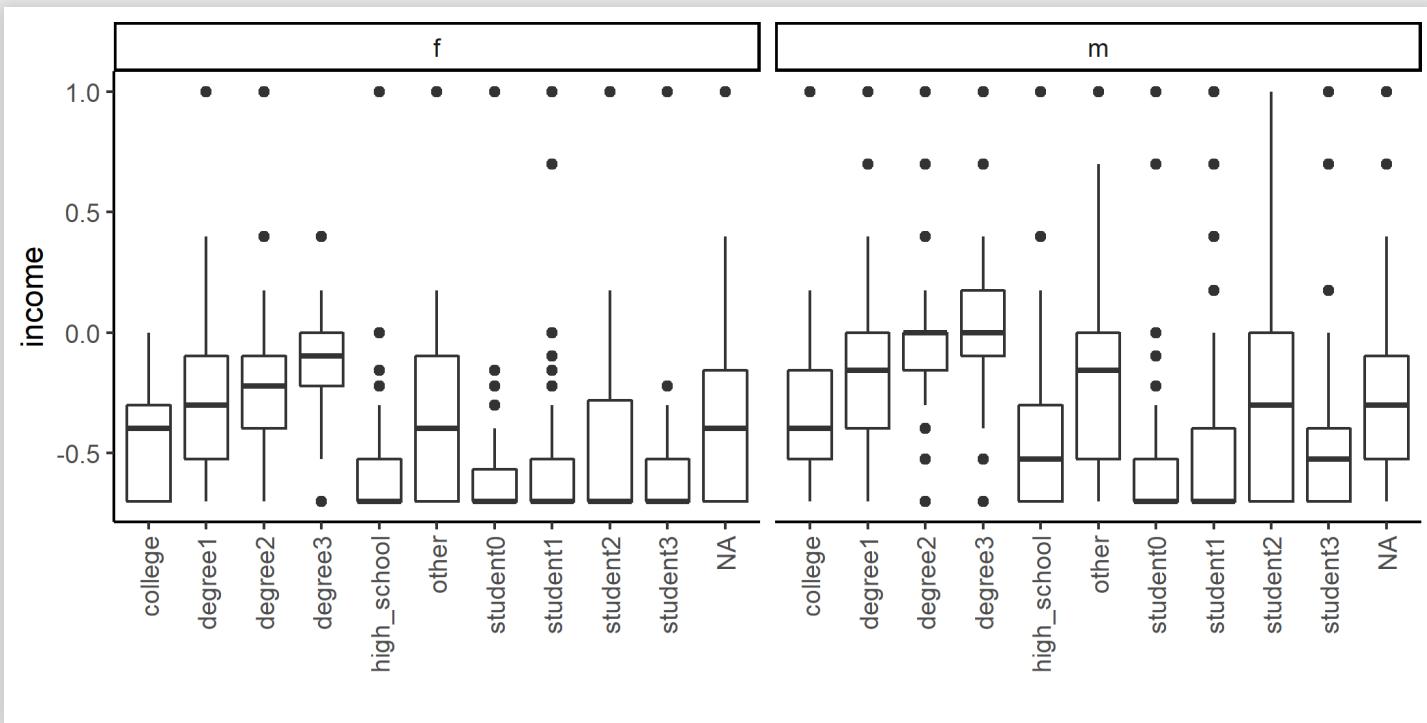
```
var_vs_income_boxplot(sign2)
```



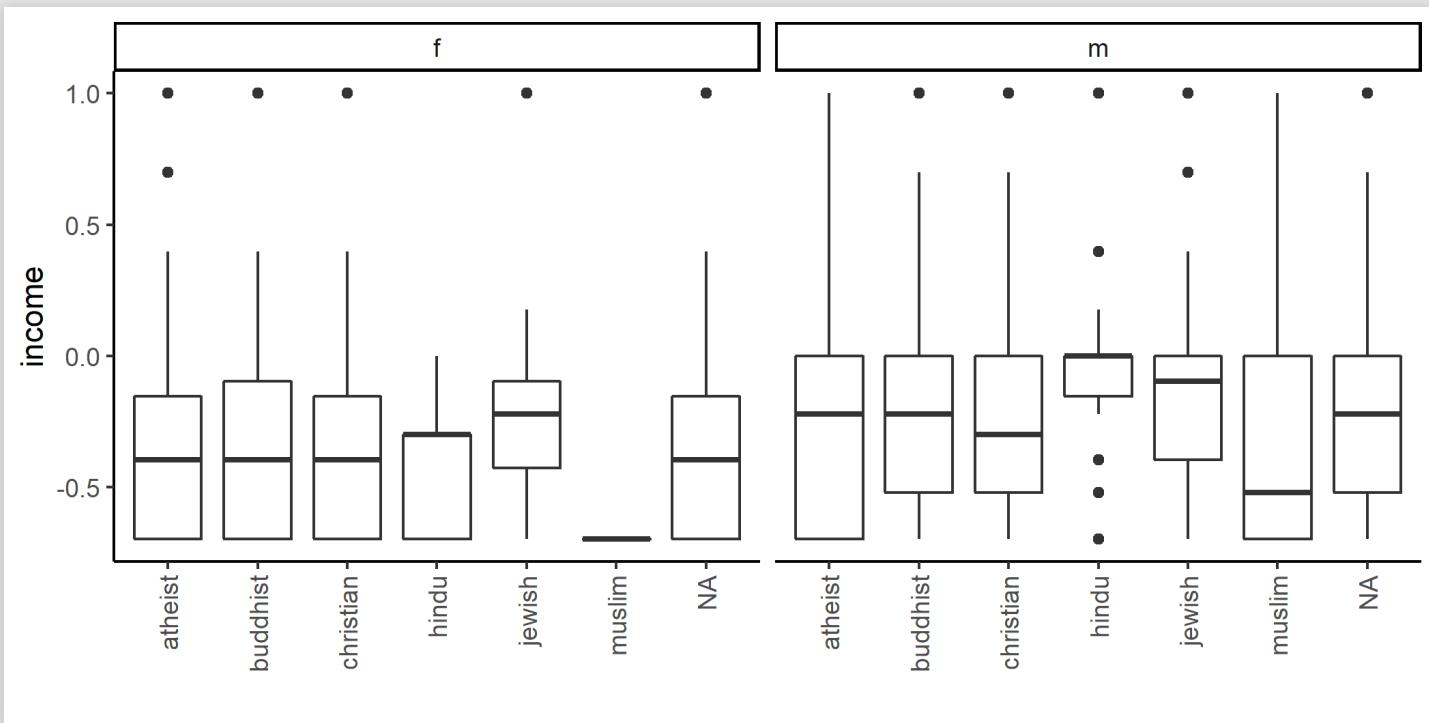
```
var_vs_income_boxplot(job3)
```



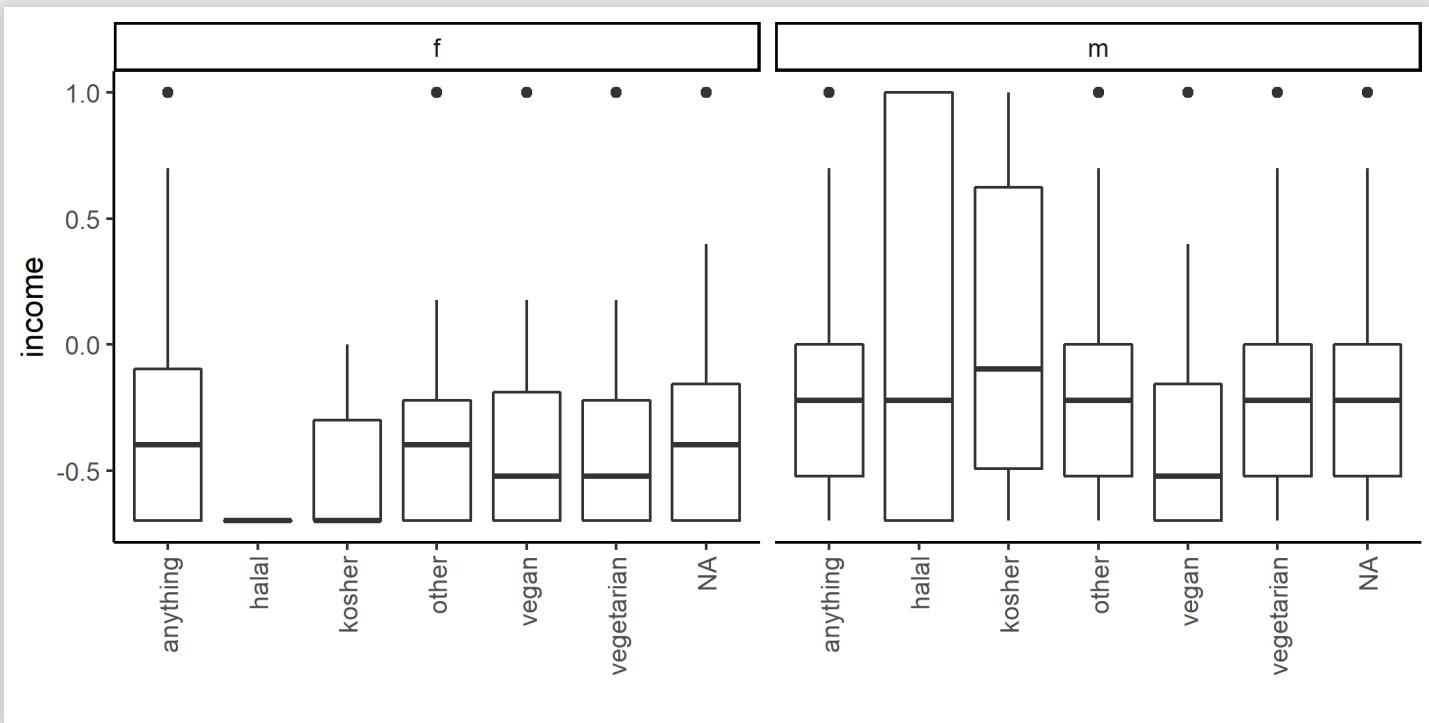
```
var_vs_income_boxplot(education2)
```




```
var_vs_income_boxplot(religion2)
```



```
var_vs_income_boxplot(diet2)
```



```
ggpairs(okcupid2_train %>%
         select(essay0_len:essay2_len, income))
```



Baseline: Linear Regression

R's `lm` function does not take NA values.

One strategy is to impute these values using a "common" value such as the median for continuous variables and mode for categorical variables. This can easily be achieved with `naniar`:

```
okcupid2_imp <- naniar::impute_median_all(okcupid2)
okcupid2_imp_train <- okcupid2_imp[train_idx, ]
okcupid2_imp_valid <- okcupid2_imp[valid_idx, ]

mod_lm <- lm(income ~ ., data = okcupid2_imp_train)
pred_lm <- predict(mod_lm, okcupid2_imp_valid)


rmse <- function(obs, pred) sqrt(mean((obs - pred)^2))

report_rmse_and_cor <- function(obs, pred) {
  RMSE <- rmse(obs, pred)
  CORR <- cor(obs, pred)
  glue("RMSE: {format(RMSE, digits = 3)}\n      CORR: {format(CORR, digits = 3)}")
}
```

```
report_rmse_and_cor(okcupid2_valid$income, pred_lm)
```

```
## RMSE: 0.352  
## CORR: 0.501
```

```
tibble(income = okcupid2_valid$income, pred = pred_lm) %>%  
  ggplot(aes(income, pred)) + geom_point() + ylim(range(okcupid2_`
```



A more intelligent strategy for imputing missing values would be to *predict* them using whatever data is not missing. This can be done quite seamlessly with the [mice](#) package:

```
# library(mice)
# mice_obj <- mice(okcupid2, m = 1, maxit = 10, seed = 42)
# okcupid2_imp_mice <- complete(mice_obj)

okcupid2_imp_mice <- read_rds("../data/okcupid2_imp_mice.rds")
okcupid2_imp_mice_train <- okcupid2_imp_mice[train_idx, ]
okcupid2_imp_mice_valid <- okcupid2_imp_mice[valid_idx, ]

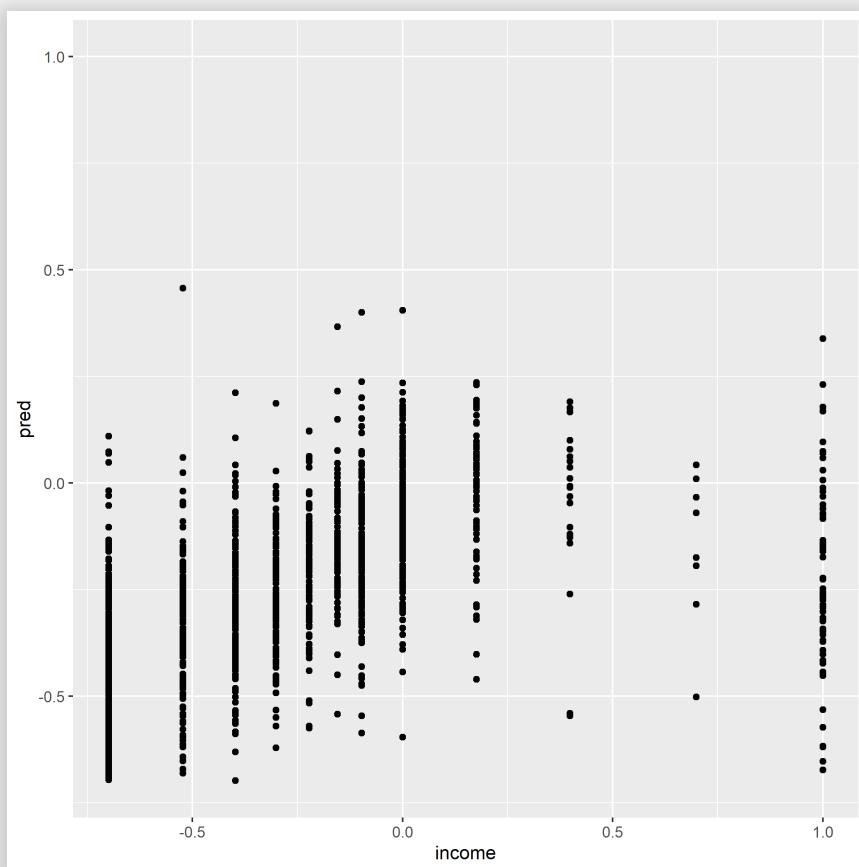
mod_lm_mice <- lm(income ~ ., data = okcupid2_imp_mice_train)
pred_lm_mice <- predict(mod_lm_mice, okcupid2_imp_mice_valid)

report_rmse_and_cor(okcupid2_valid$income, pred_lm_mice)

## RMSE: 0.351
## CORR: 0.504
```

 Can you think of other imputation strategies?

```
tibble(income = okcupid2_valid$income, pred = pred_lm_mice) %>%  
  ggplot(aes(income, pred)) +  
  geom_point() +  
  ylim(range(okcupid2_valid$income))
```



Baseline: Ridge Regression

```
library(glmnet)

okcupid2_imp_mat_train <- model.matrix(~ ., okcupid2_imp_mice_train)
okcupid2_imp_mat_valid <- model.matrix(~ ., okcupid2_imp_mice_valid)

ridge_cv <- cv.glmnet(x = okcupid2_imp_mat_train,
                      y = okcupid2_train$income, alpha = 0)

best_lambda <- ridge_cv$lambda.min

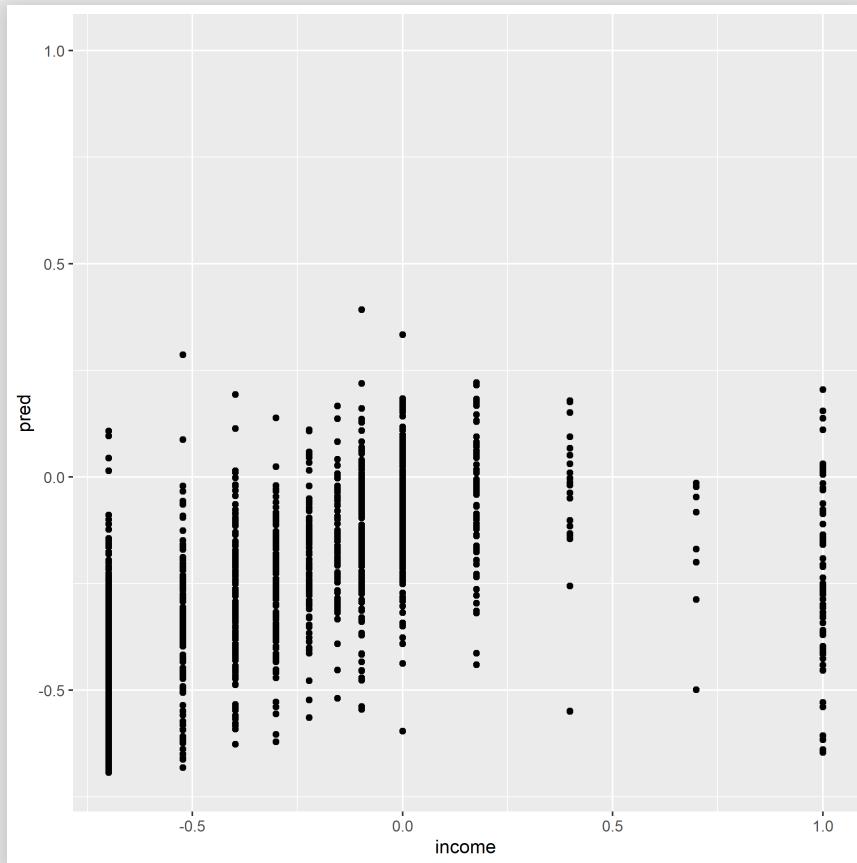
mod_lm_ridge <- glmnet(x = okcupid2_imp_mat_train,
                        y = okcupid2_train$income, alpha = 0,
                        lambda = best_lambda)

pred_lm_ridge <- predict(mod_lm_ridge, okcupid2_imp_mat_valid)

report_rmse_and_cor(okcupid2_valid$income, pred_lm_ridge)

## RMSE: 0.351
## CORR: 0.505
```


```
tibble(income = okcupid2_valid$income, pred = pred_lm_ridge) %>%
  ggplot(aes(income, pred)) +
  geom_point() +
  ylim(range(okcupid2_valid$income))
```



Baseline: Lasso Regression

```
lasso_cv <- cv.glmnet(x = okupid2_imp_mat_train,  
                      y = okupid2_imp_train$income, alpha = 1)  
  
best_lambda <- lasso_cv$lambda.min  
  
mod_lm_lasso <- glmnet(x = okupid2_imp_mat_train,  
                        y = okupid2_train$income, alpha = 1,  
                        lambda = best_lambda)  
  
pred_lm_lasso <- predict(mod_lm_lasso, okupid2_imp_mat_valid)  
  
report_rmse_and_cor(okupid2_valid$income, pred_lm_lasso)  
  
## RMSE: 0.351  
## CORR: 0.506
```

```
tibble(income = okcupid2_valid$income, pred = pred_lm_lasso) %>%
  ggplot(aes(income, pred)) +
  geom_point() +
  ylim(range(okcupid2_valid$income))
```



End of Detour

APPLICATIONS



OF DATA SCIENCE

The CART (Regression)

APPLICATIONS



OF DATA SCIENCE

The OG CART

1. Find the predictor to (binary) split on and the value of the split, by SSE criterion
2. For each resulting node if $n_{node} > 20$ go to 1
3. Once full tree has been grown, perform pruning using the *cost-complexity parameter* c_p and the SSE_{c_p} criterion
4. Predict the average value at each terminal node
 - For each split *surrogate splits* are saved for future NAs
 - An alternative criterion for complexity could be tree maximum depth (sklearn)
 - One can also reduce all tree's unique paths to a set of rules
 - Variables can be ranked by "importance"



What if the best split isn't binary?

The SSE criterion - continuous predictor

- y is the continuous dependent variable
- a continuous predictor v is nominated for splitting the current node
- with splitting value l
- such that S_1 is the set of observations for which $v_i \leq l$
- and S_2 is the set of observations for which $v_i > l$
- \bar{y}_1 is the average of y in set S_1

$$SSE = \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2$$

For example if `age` is candidate in splitting `income`:

```

library(patchwork)

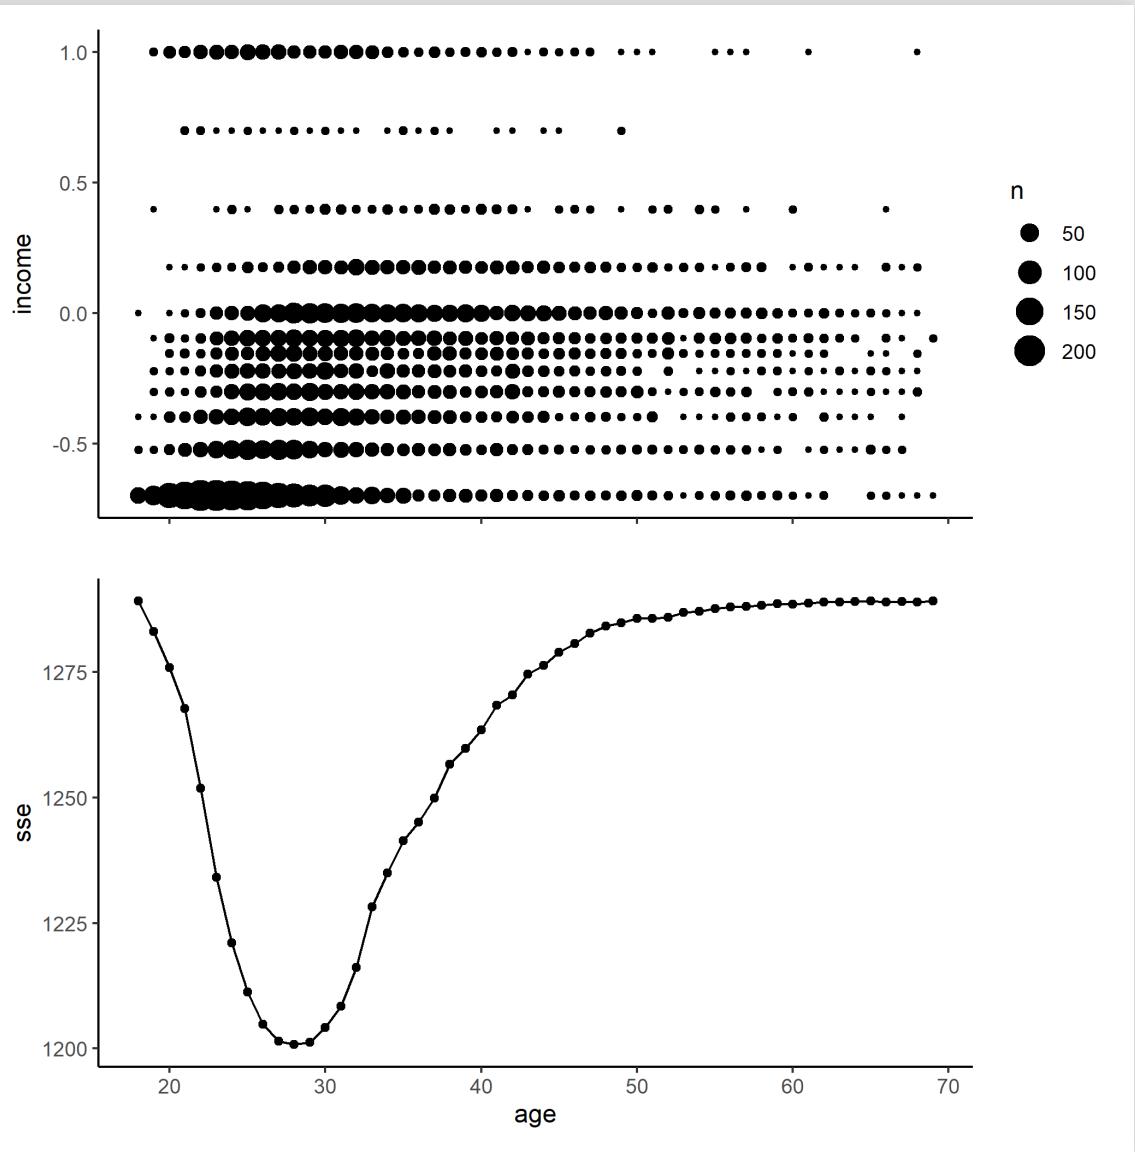
sse <- function(l, df, v) {
  income_above <- df %>% filter({{v}} >= l) %>% pull(income)
  income_below <- df %>% filter({{v}} < l) %>% pull(income)
  sse_above <- sum((income_above - mean(income_above))^2)
  sse_below <- sum((income_below - mean(income_below))^2)
  return(sse_above + sse_below)
}
age <- seq(18, 69, 1)
sse_age <- map_dbl(age, sse, df = okcupid2_train, v = age)

p1 <- okcupid2_train %>%
  count(age, income) %>%
  ggplot(aes(age, income)) +
  geom_point(aes(size = n)) +
  theme_classic() +
  labs(x = "") +
  theme(axis.text.x = element_blank())

p2 <- tibble(age = age, sse = sse_age) %>%
  ggplot(aes(age, sse)) +
  geom_line() +
  geom_point() +
  theme_classic()

p1 / p2

```



The *SSE* criterion - categorical predictor

💡 What could be an issue with a categorical variable with many levels?

- a categorical predictor l is nominated for splitting the current node
 - option 1: use dummy variables, turning each level into a 2-level 0/1 category variable
 - option 2: order levels by some criterion like \bar{y}_j and treat l as continuous from here on

For example if `job3` is candidate in splitting `income`:

```

mean_income_vs_job <- okcupid2_train %>%
  group_by(job3) %>%
  summarise(mean_income = mean(income)) %>%
  arrange(mean_income)
jobs_levels_sorted <- as.character(mean_income_vs_job$job3)
okcupid2_train_job_sorted <- okcupid2_train %>%
  mutate(job_ordered = fct_relevel(job3, jobs_levels_sorted),
         job_ordered_n = as.numeric(job_ordered))

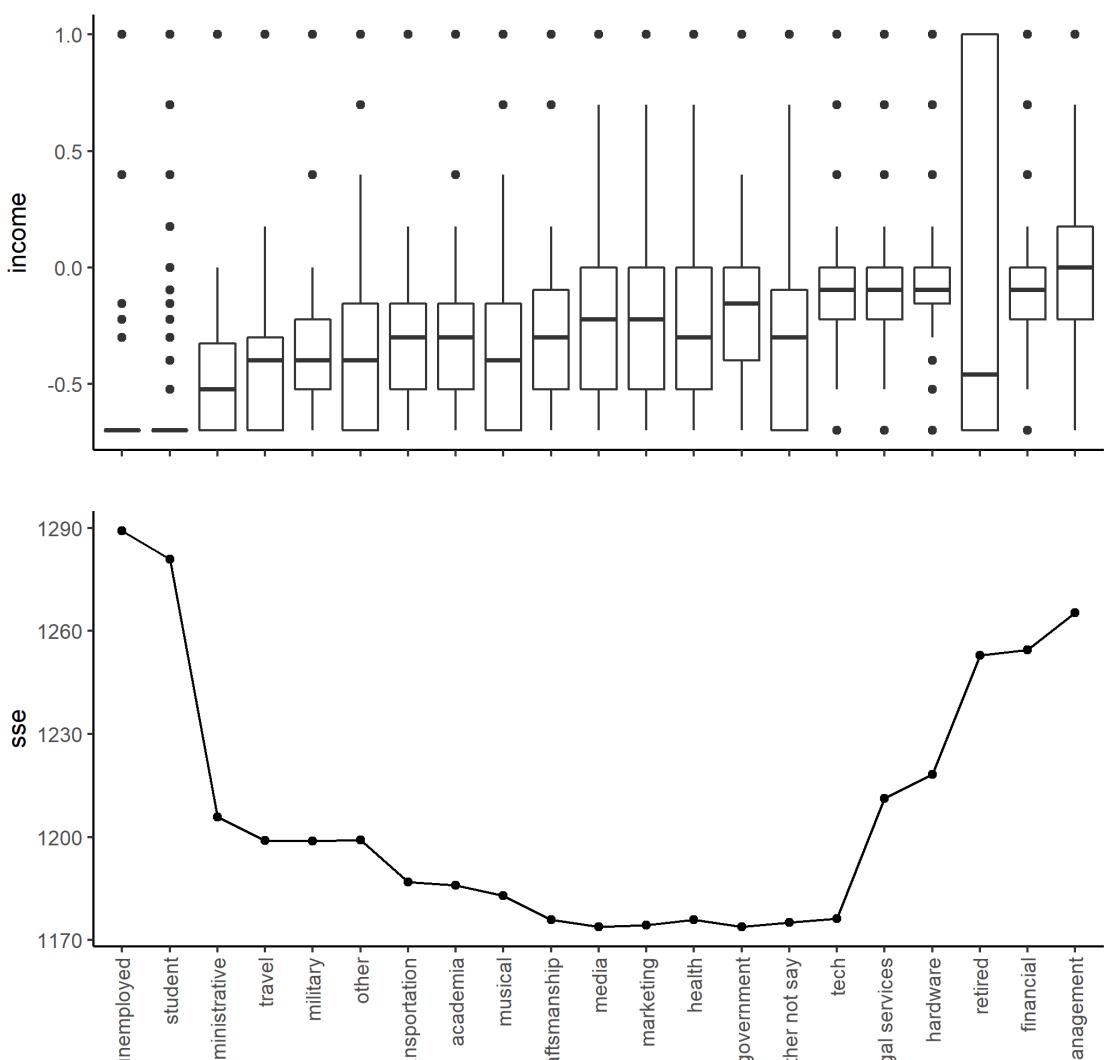
job_rank <- seq(1, length(jobs_levels_sorted), 1)
sse_job <- map_dbl(job_rank, sse, df = okcupid2_train_job_sorted,

p1 <- okcupid2_train_job_sorted %>%
  ggplot(aes(job_ordered, income)) +
  geom_boxplot() +
  theme_classic() + labs(x = "") +
  theme(axis.text.x = element_blank())

p2 <- tibble(job = factor(jobs_levels_sorted, levels = jobs_levels_sorted),
             sse, group = 1)) +
  geom_line() +
  geom_point() +
  theme_classic() +
  labs(x = "") +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust = 0))

p1 / p2

```



Pruning: The SSE_{c_p} criterion

The "right-sized" tree should not be too deep to avoid overfitting.

Once the full tree is grown, we check for each split:

$$SSE_{c_p} = SSE_{tot} + c_p \cdot \#\text{terminal nodes}$$

Where c_p is a penalty or regularization parameter usually chosen by cross validation.

And we choose the smallest pruned tree with the minimum SSE_{c_p} .

CART with rpart

Growing the full tree (for a numeric (y) rpart will guess this is Regression):

```
library(rpart)  
  
mod_tree <- rpart(income ~ ., data = okcupid2_train)
```

Pruning with some c_p :

```
mod_tree <- prune(mod_tree, cp = 0.05)
```

⚠ There is a c_p parameter you can pass to rpart while training like so:

```
rpart(income ~ ., data = okcupid2_train, control =  
rpart.control(cp = 0.05))
```

But this will only make rpart consider this c_p at each split as a minimum criterion *while growing the unpruned tree*. In fact the default of this parameter is 0.01!

So, in order to train a true unpruned tree you would need to pass $c_p = 0$:

```
mod_tree <- rpart(income ~ ., data = okcupid2_train,  
control = rpart.control(cp = 0))
```

Now, `rpart` by default will perform 10-fold Cross Validation on each split, resulting in this `cptable`:

```
head(mod_tree$cptable)
```

```
##          CP nsplit rel_error xerror      xstd  
## 1 0.113597498    0 1.0000000 1.0002522 0.02238510  
## 2 0.048179174    1 0.8864025 0.8905350 0.02282679  
## 3 0.013140594    2 0.8382233 0.8454403 0.02367885  
## 4 0.010043056    3 0.8250827 0.8383955 0.02392361  
## 5 0.009147817    4 0.8150397 0.8301138 0.02419424  
## 6 0.005634308    5 0.8058919 0.8242994 0.02431132
```

The nature of the `xerror` is unclear from the [docs](#) (SSE/n ?) except that it is relative to the error in the root node.

You can either use it like so:

```
best_cp <- mod_tree$cptable[which.min(mod_tree$cptable[, "xerror"])]
mod_tree <- prune(mod_tree, cp = best_cp)
```

Or you can perform CV on your own, passing at each stage a parameter for `rpart` to not perform CV:

```
mod_tree <- rpart(income ~ ., data = okcupid2_train,
control = rpart.control(cp = 0), xval = 1)
```

Let's tune c_p for our data using a 5-fold (manual) Cross Validation.
The criterion to maximize would be RMSE.

```
n_cv <- 5; cp_seq <- seq(0, 0.02, 0.001)

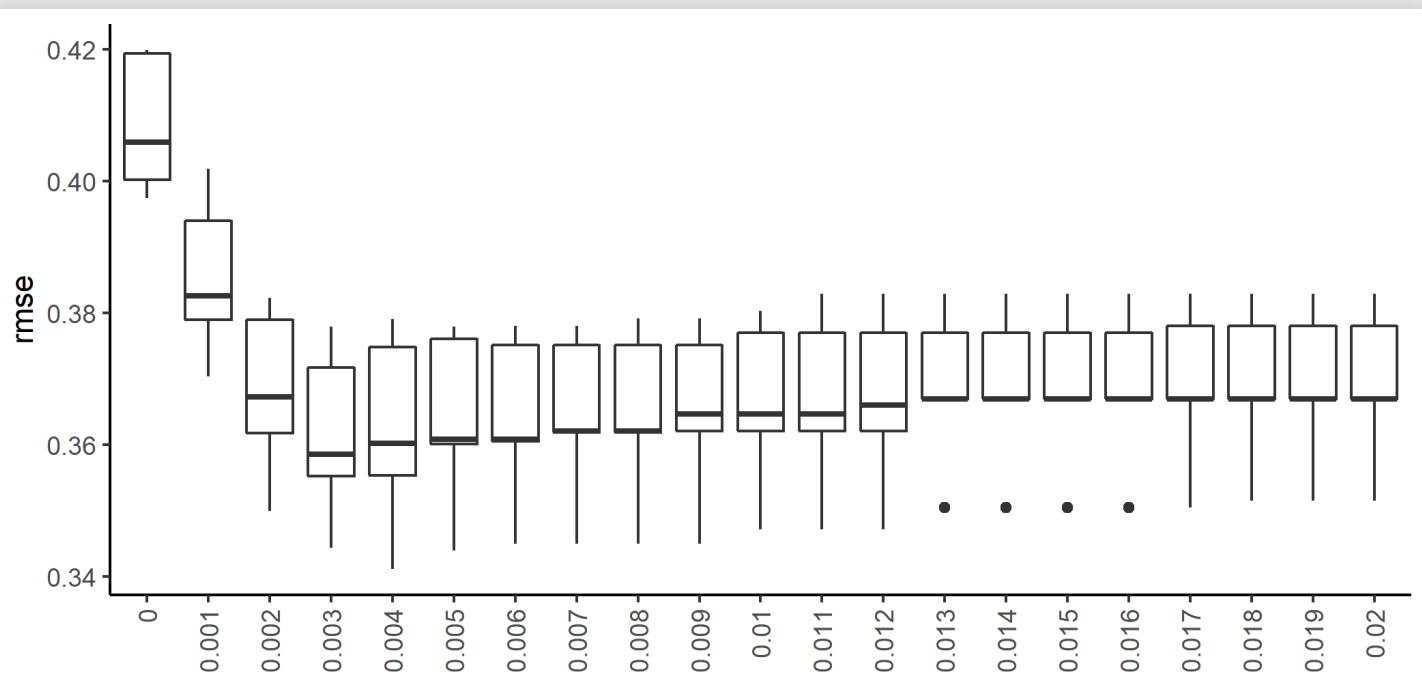
okupid2_train_val <- okupid2_train %>%
  mutate(val = sample(1:n_cv, n(), replace = TRUE))

get_validation_set_rmse <- function(i, .cp) {
  ok_tr <- okupid2_train_val %>% filter(val != i) %>% select(-val)
  ok_val <- okupid2_train_val %>% filter(val == i) %>% select(-val)
  mod <- rpart(income ~ ., data = ok_tr,
               control = rpart.control(cp = 0, xval = 1))
  mod <- prune(mod, cp = .cp)
  pred <- predict(mod, ok_val)
  rmse(ok_val$income, pred)
}

get_cv_rmse <- function(.cp) {
  tibble(cp = rep(.cp, n_cv),
        rmse = map_dbl(1:n_cv, get_validation_set_rmse, .cp = .cp)
}
```

```
cv_table <- map_dfr(cp_seq, get_cv_rmse)

cv_table %>%
  mutate(cp = factor(cp)) %>%
  ggplot(aes(cp, rmse)) +
  geom_boxplot() +
  theme_classic() +
  labs(x = "") +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust =
```



Training on the entire training set:

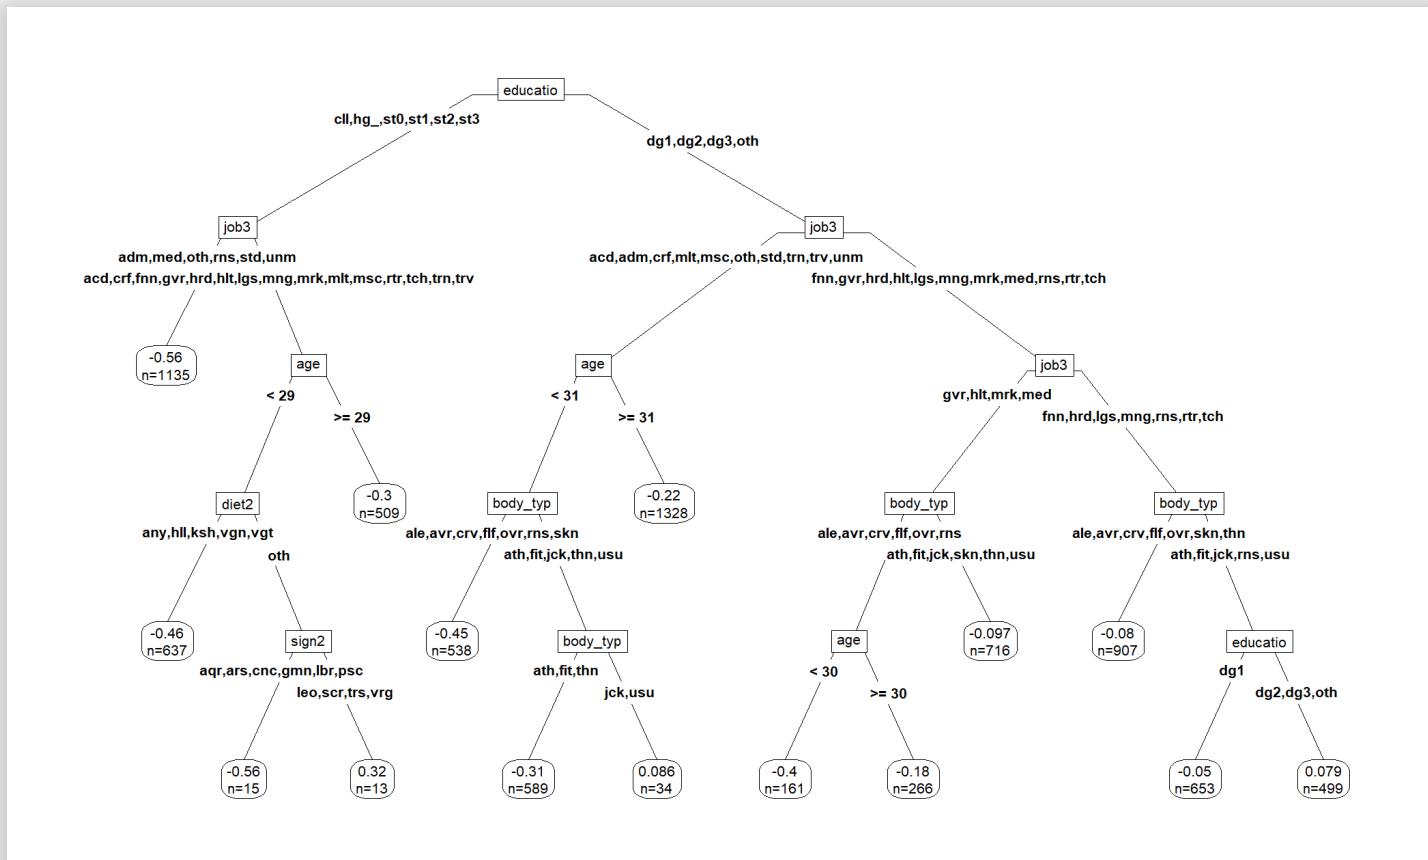
```
mod_tree_na <- rpart(income ~ ., data = okcupid2_train,  
                      control = rpart.control(cp = 0, xval = 1))  
mod_tree_na <- prune(mod_tree_na, cp = 0.003)
```

You can plot the tree using `rpart`, but...

```
plot(mod_tree_na)  
text(mod_tree_na, pretty = 1, use.n = TRUE)
```

The plotting function in the `rpart.plot` package is slightly nicer:

```
library(rpart.plot)
prp(mod_tree_na, type = 5, extra = 1)
```

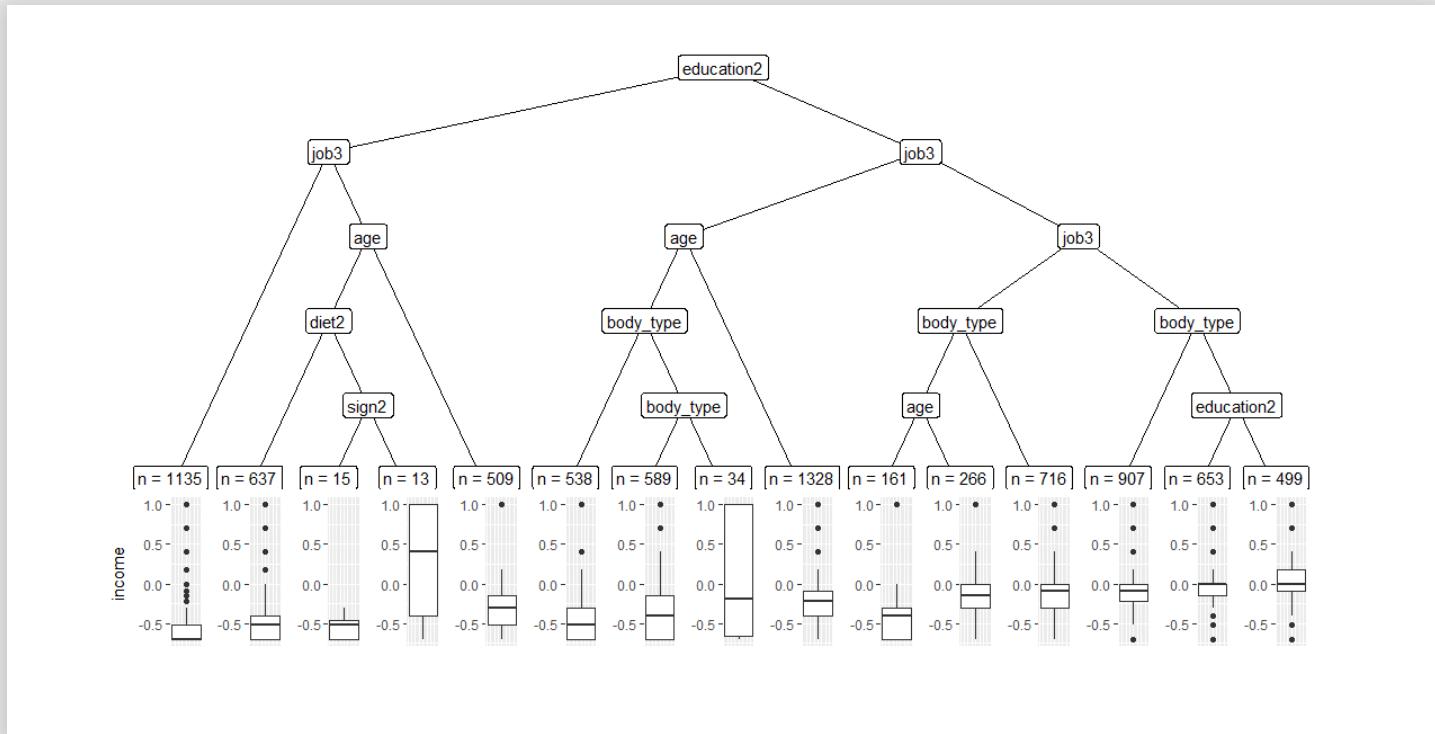


You can go fancy and use partykit and/or ggparty:

```
library(ggparty)

party_obj <- as.party(mod_tree_na)

ggparty(party_obj) +
  geom_edge() +
  # geom_edge_label() +
  geom_node_label(aes(label = splitvar), ids = "inner") +
  geom_node_label(aes(label = str_c("n = ", nodesize)),
                  ids = "terminal", nudge_y = 0.02) +
  geom_node_plot(gglist = list(geom_boxplot(aes(y = income)),
                               theme(axis.text.x=element_blank(),
                                     axis.ticks.x=element_blank,
                                     shared_axis_labels=TRUE))
```



Or you can just print the tree and see what is going on:

```
print(party_obj)
```

```
Fitted party:
[1] root
[2] education2 in college, high_school, student0, student1, student2, student3
| [3] job3 in administrative, media, other, rather not say, student, unemployed: -0.555 (n = 1135, err = 126.9)
| [4] job3 in academia, craftsmanship, financial, government, hardware, health, legal services, management, marketing, military
| [5] age < 28.5
|   [6] diet2 in anything, halal, kosher, vegan, vegetarian: -0.460 (n = 637, err = 84.2)
|   [7] diet2 in other
|     [8] sign2 in aquarius, aries, cancer, gemini, libra, pisces: -0.559 (n = 15, err = 0.3)
|     [9] sign2 in leo, scorpio, taurus, virgo: 0.316 (n = 13, err = 5.9)
|   [10] age >= 28.5: -0.301 (n = 509, err = 53.0)
[11] education2 in degree1, degree2, degree3, other
[12] job3 in academia, administrative, craftsmanship, military, musical, other, student, transportation, travel, unemployed
| [13] age < 30.5
|   [14] body_type in a little extra, average, curvy, full figured, overweight, rather not say, skinny: -0.450 (n = 538,
|   [15] body_type in athletic, fit, jacked, thin, used up
|     [16] body_type in athletic, fit, thin: -0.314 (n = 589, err = 141.6)
|     [17] body_type in jacked, used up: 0.086 (n = 34, err = 16.1)
|   [18] age >= 30.5: -0.221 (n = 1328, err = 171.8)
[19] job3 in financial, government, hardware, health, legal services, management, marketing, media, rather not say, retired,
[20] job3 in government, health, marketing, media
| [21] body_type in a little extra, average, curvy, full figured, overweight, rather not say
|   [22] age < 29.5: -0.402 (n = 161, err = 13.7)
|   [23] age >= 29.5: -0.178 (n = 266, err = 17.3)
|     [24] body_type in athletic, fit, jacked, skinny, thin, used up: -0.097 (n = 716, err = 95.6)
[25] job3 in financial, hardware, legal services, management, rather not say, retired, tech
| [26] body_type in a little extra, average, curvy, full figured, overweight, skinny, thin: -0.080 (n = 907, err = 83.
|   [27] body_type in athletic, fit, jacked, rather not say, used up
|     [28] education2 in degree1: -0.050 (n = 653, err = 48.7)
|     [29] education2 in degree2, degree3, other: 0.079 (n = 499, err = 51.3)

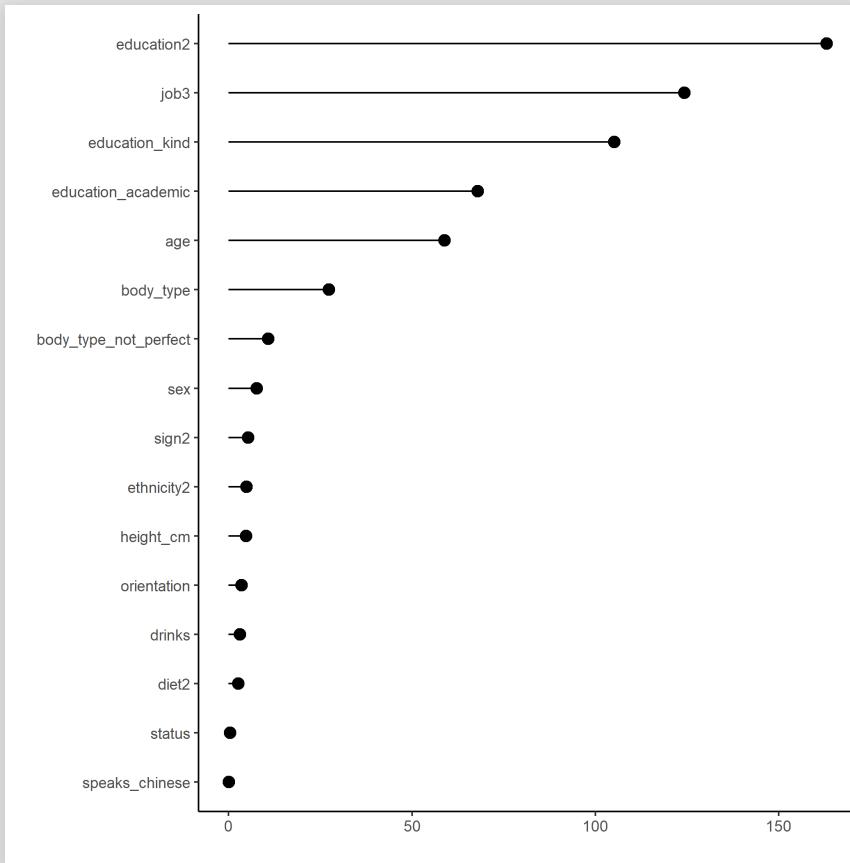
Number of inner nodes: 14
Number of terminal nodes: 15
```

Variables Importance

Summing the reduction in SSE for each split variable, we can get a measure of importance.

Unfortunately in `rpart` the "potential" reduction in (SSE) for surrogate splits is also summed. You can either ignore or retrain with only the variables chosen by the model.

```
enframe(mod_tree_na$variable.importance) %>%
  arrange(value) %>%
  mutate(variable = as_factor(name)) %>%
  ggplot(aes(variable, value)) +
  geom_segment(aes(x = variable, xend = variable,
                    y = 0, yend = value)) +
  geom_point(size = 3) +
  theme_classic() +
  coord_flip() +
  labs(x = "", y = "")
```



💡 How would this profile look for a couple of very correlated predictors?

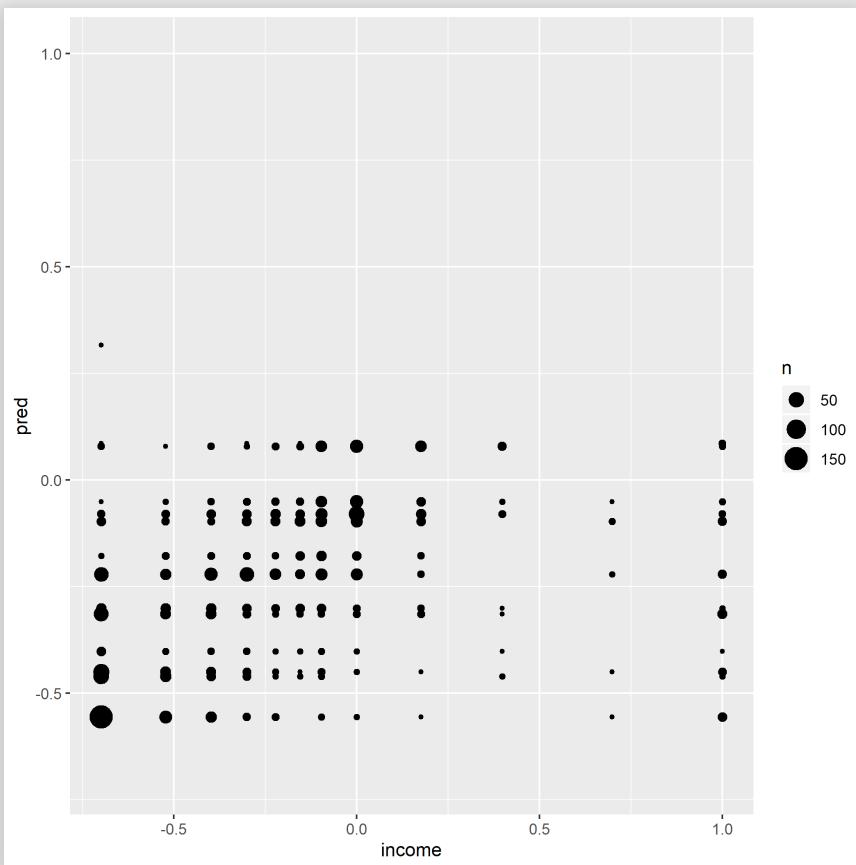
Prediction

```
pred_tree_na <- predict(mod_tree_na, okcupid2_valid)  
report_rmse_and_cor(okcupid2_valid$income, pred_tree_na)  
  
## RMSE: 0.363  
## CORR: 0.449
```

As expected, far from impressive. Let's try using the imputed NA data:

```
mod_tree_imp <- rpart(income ~ ., data = okcupid2_imp_mice_train,  
                      control = rpart.control(cp = 0, xval = 1))  
mod_tree_imp <- prune(mod_tree_imp, cp = 0.003)  
pred_tree_imp <- predict(mod_tree_imp, okcupid2_imp_mice_valid)  
  
report_rmse_and_cor(okcupid2_valid$income, pred_tree_imp)  
  
## RMSE: 0.369  
## CORR: 0.424
```

```
tibble(income = okcupid2_valid$income, pred = pred_tree_na) %>%
  count(income, pred) %>%
  ggplot(aes(income, pred)) +
  geom_point(aes(size = n)) +
  ylim(range(okcupid2_valid$income))
```

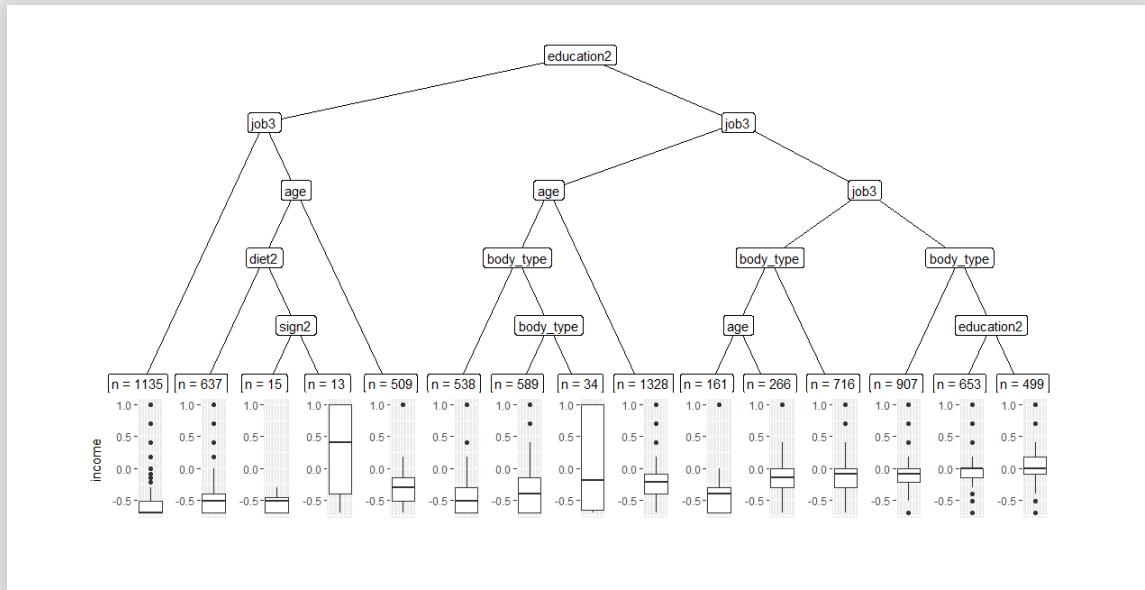


The Others

APPLICATIONS



OF DATA SCIENCE



Looking at this one might wonder:

- Since when do we just split, how about a statistical test?
- Predicting a single value for each terminal node? Is that the best we can do?
- Many variables repeat in the same path, can we compact paths into simple rules?

Conditional Inference Trees

[Hothorn et. al. \(2006\)](#) perform a statistical hypothesis test for each variable to decide whether it is related to y (including multiple comparisons adjustment).

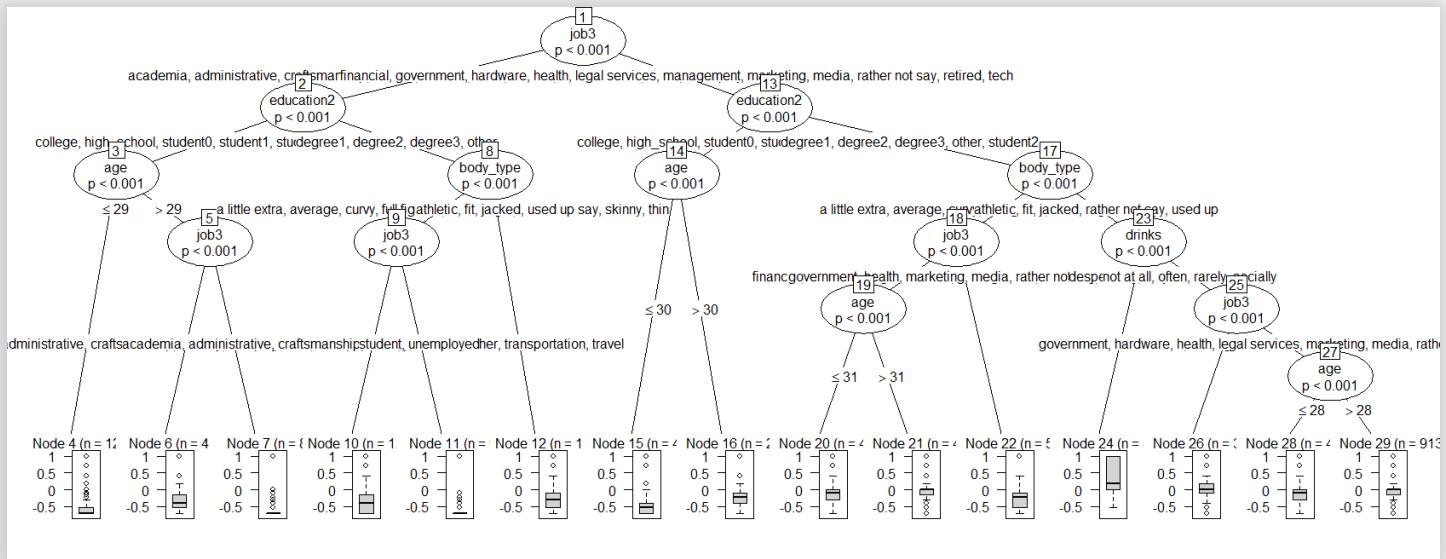
If no variable passes test - stop. No pruning necessary. If a few variable pass - choose the most related, e.g. using a p-value.

```
library(partykit)

mod_ctree_na <- ctree(income ~ ., data = okcupid2_train, alpha = 0)
pred_ctree_na <- predict(mod_ctree_na, okcupid2_valid)
report_rmse_and_cor(okcupid2_valid$income, pred_ctree_na)

## RMSE: 0.363
## CORR: 0.449
```

```
plot(mod_ctree_na)
```



Model Trees

[Quinlan \(1992\)](#) and later [Wang and Witten \(1997\)](#) made a few changes to the beloved CART:

- instead of single values at terminal nodes, M5 trees have *linear models*
- sort of a piecewise linear regression only the "pieces" can come from many variables
- can extrapolate to never before seen values (not necessarily good)

You can find an implementation in R in the `RWeka` package, function `M5P()`.

[Zeileis et. al. \(2008\)](#) approach is slightly simpler to grasp. It allows you to specify variables for splitting and variables for regression.

See `lmtree()` and `glmtree()` functions in the `partykit`

```

mod_lmtree_na <- lmtree(income ~ age | sex + job3 + education2, data = okupid2)

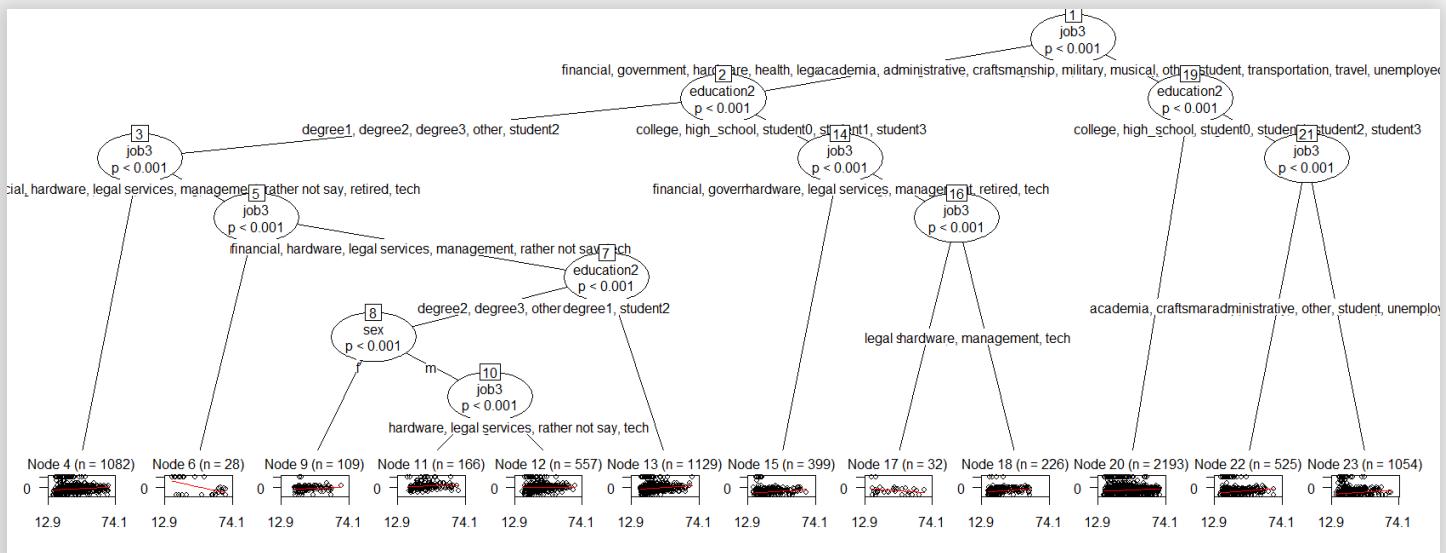
pred_lmtree_na <- predict(mod_lmtree_na, okupid2_valid)

report_rmse_and_cor(okupid2_valid$income, pred_lmtree_na)

plot(mod_lmtree_na)

```

RMSE: 0.366
CORR: 0.437



Rule-Based Trees

Holmes et. a. (1993) general approach:

1. Build a model tree (e.g. M5)
2. For each path compute the coverage (= % of relevant observations to that path)
3. Compact the path with largest coverage into a rule
4. Remove previous relevant observations and go back to 1 until all observations are accounted for.

```
library(RWeka)

# limiting predictors just for a nice print!
mod_m5rules_na <- M5Rules(income ~ ., data = okcupid2_train[, c("i
pred_m5rules_na <- predict(mod_m5rules_na, okcupid2_valid)
report_rmse_and_cor(okcupid2_valid$income, pred_m5rules_na)

## RMSE: 0.374
## CORR: 0.392
```

```
print(mod_m5rules_na)
```

```
M5 pruned model rules
(using smoothed linear models) :
Number of Rules : 2

Rule: 1
IF
    education2=student2,degree1,other,degree2,degree3 > 0.5
    age > 30.5
THEN
income =
    -0.0015 * age
    + 0.1123 * sex=m
    + 0.0003 * education2=student1,high_school,student3,college,student2,degree1,other,degree2,degree3
    + 0.0002 * education2=high_school,student3,college,student2,degree1,other,degree2,degree3
    + 0.0002 * education2=student2,degree1,other,degree2,degree3
    + 0.185 * education2=degree1,other,degree2,degree3
    + 0.0003 * education2=other,degree2,degree3
    + 0.0672 * education2=degree2,degree3
    + 0.0755 * education2=degree3
    - 0.3345 [2977/81.357%]

Rule: 2
income =
    0.009 * age
    + 0.1196 * sex=m
    + 0.0927 * education2=student1,high_school,student3,college,student2,degree1,other,degree2,degree3
    - 0.0384 * education2=high_school,student3,college,student2,degree1,other,degree2,degree3
    + 0.0505 * education2=student3,college,student2,degree1,other,degree2,degree3
    + 0.212 * education2=student2,degree1,other,degree2,degree3
    + 0.111 * education2=degree2,degree3
    - 0.8955 [4523/93.176%]
```

Bagged Trees

APPLICATIONS



Of DATA SCIENCE

Bagging: The Gist


For $j = 1$ to m do

- Generate a bootstrap sample of the original data
- Train an unpruned tree model on this sample

End.

- Predict average through all m trees (or %vote for each class in classification)
- What you get: better prediction, OOB error estimates
- What you lose: interpretability, speed (but bagging can be paralleled)

Why does Bagging work?



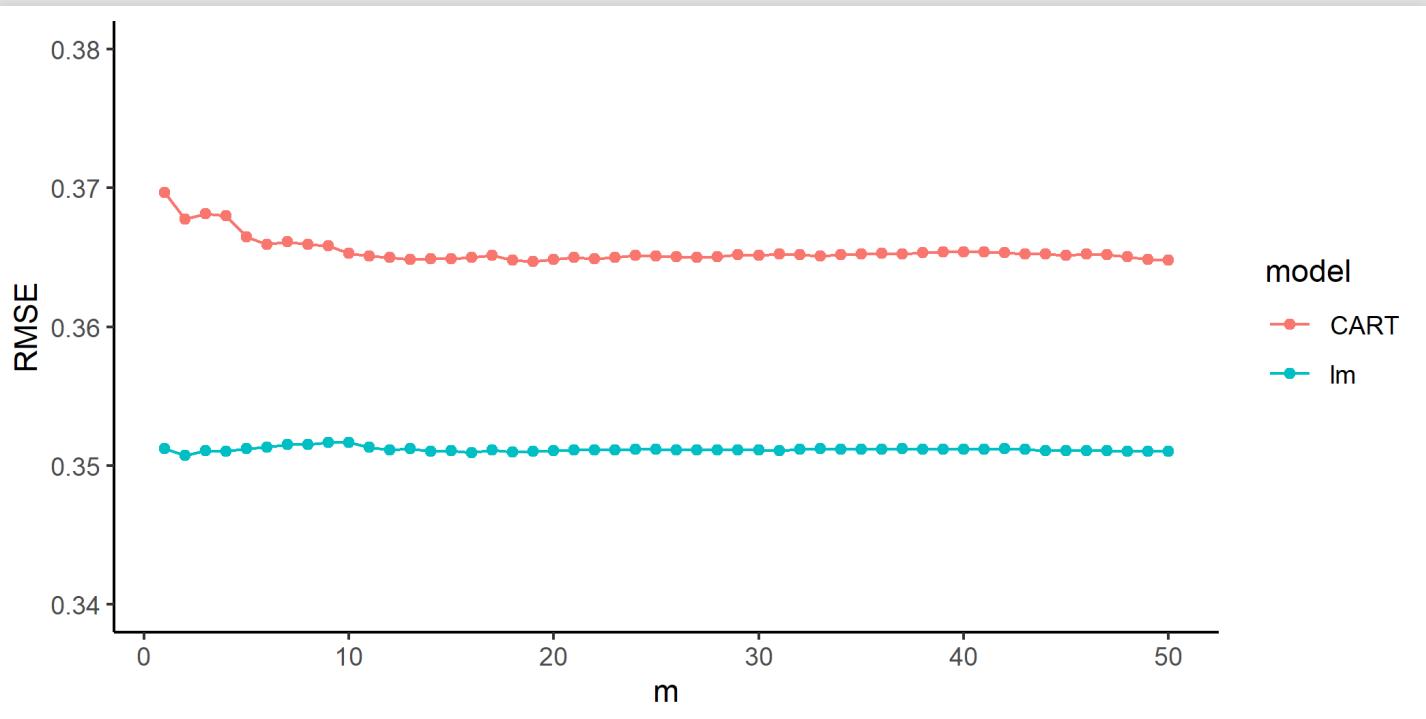
Which models are likely to benefit from bagging?

```
train_lm <- function(i, .n) {  
  samp <- sample(1:n, .n, replace = TRUE)  
  lm(income ~ ., data = okcupid2_imp_mice_train[samp, ])  
}  
  
train_tree <- function(i, .n) {  
  samp <- sample(1:n, .n, replace = TRUE)  
  rpart(income ~ ., data = okcupid2_imp_mice_train[samp, ])  
}  
  
predict_mod <- function(mod) {  
  tibble(pred = predict(mod, okcupid2_imp_mice_valid))  
}  
  
rmse_m_models <- function(m, .preds) {  
  rmse(okcupid2_imp_mice_valid$income, rowMeans(.preds[, 1:m]))  
}  
  
n <- nrow(okcupid2_imp_mice_train)  
m_max <- 50  
  
mod_lms <- map(1:m_max, train_lm, .n = n)  
mod_trees <- map(1:m_max, train_tree, .n = n)  
  
preds_lm <- map_dfc(mod_lms, predict_mod)  
preds_tree <- map_dfc(mod_trees, predict_mod)
```

```

rmse_lm <- map_dbl(1:m_max, rmse_m_models, .preds = preds_lm)
rmse_tree <- map_dbl(1:m_max, rmse_m_models, .preds = preds_tree)
tibble(m = rep(1:m_max, 2),
       model = c(rep("lm", m_max), rep("CART", m_max)),
       RMSE = c(rmse_lm, rmse_tree)) %>%
  ggplot(aes(m, RMSE, color = model)) +
  ylim(c(0.34, 0.38)) +
  geom_line() + geom_point() +
  theme_classic()

```



Bagging with ipred

```
library(ipred)

mod_bag_na <- bagging(income ~ ., data = okcupid2_train, nbagg = 5)
pred_bag_na <- predict(mod_bag_na, okcupid2_valid)
report_rmse_and_cor(okcupid2_valid$income, pred_bag_na)

## RMSE: 0.366
## CORR: 0.438

mod_bag_mice <- bagging(income ~ ., data = okcupid2_imp_mice_train,
                          nbagg = 5)
pred_bag_mice <- predict(mod_bag_na, okcupid2_imp_mice_valid)
report_rmse_and_cor(okcupid2_valid$income, pred_bag_mice)

## RMSE: 0.364
## CORR: 0.444
```

Random Forests

APPLICATIONS



OF DATA SCIENCE


Random Forests: The Gist

For $j = 1$ to m do

- Generate a bootstrap sample of the original data
- Train an unpruned tree model on this sample
- For each splot:
 - Randomly select m_{try} predictors
 - Select best predictor for split in this subset only

End.

What does RF add on top of Bagging?



RF with randomForest


For some reason this implementation of RF won't accept missing values:

```
library(randomForest)

mod_rf_mice <- randomForest(income ~ ., data = okcupid2_imp_mice_t
                           mtry = 10, ntree = 50, importance = TRUE)
pred_rf_mice <- predict(mod_rf_mice, okcupid2_imp_mice_valid)
report_rmse_and_cor(okcupid2_valid$income, pred_rf_mice)

## RMSE: 0.354
## CORR: 0.491
```

For variable importance check out the `importance` field of the RF object:

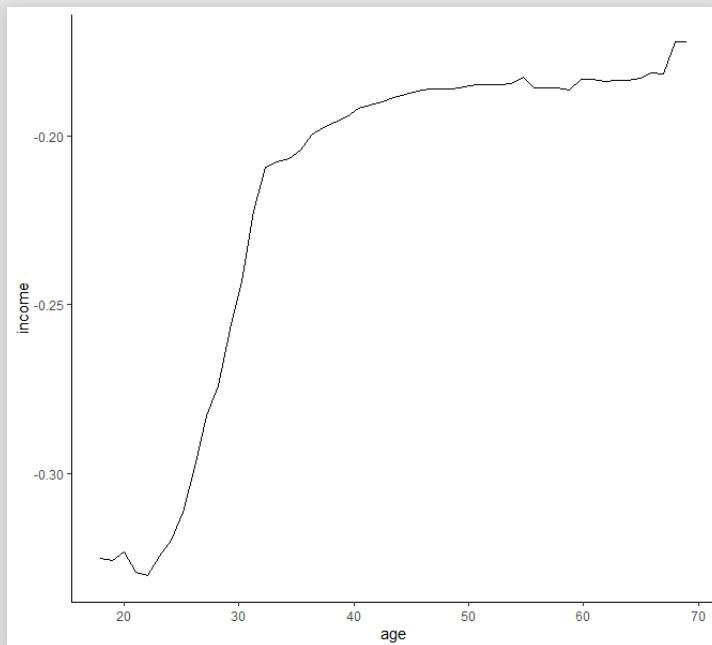


💡 What is troubling about variable importance?

Partial Dependency Plots

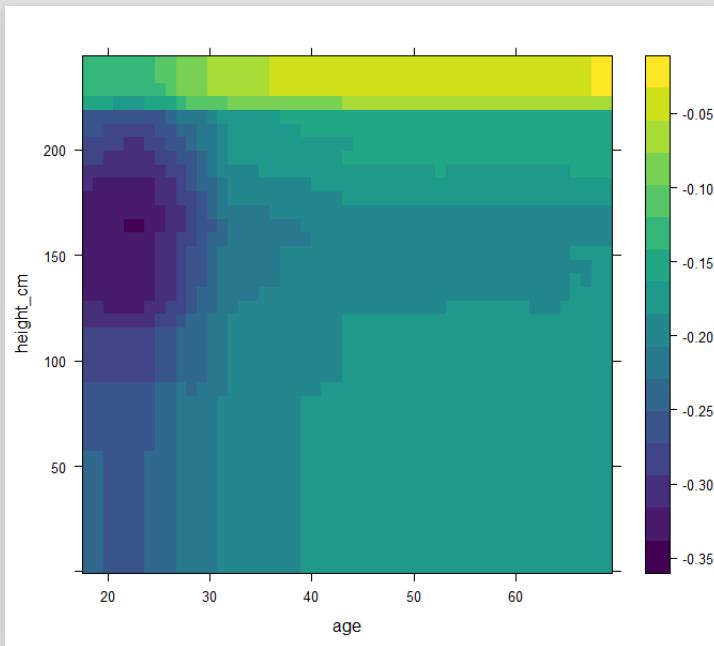
1. Select a subset X_s of variables (usually 1-2) you wish to know the relation to the dependent variable y
2. For each possible permutation of X_s run the model on the original data imputing this specific permutation instead of the values of X_s
3. Average predictions of y
4. Plot y vs. X_s

```
partial(mod_rf_mice, pred.var = "age",
       plot = TRUE, plot.engine = "ggplot2") +
  labs(y = "income") +
  theme_classic()
```



Let's see this for age and height_cm together (takes some time!):

```
partial(mod_rf_mice, pred.var = c("age", "height_cm"), plot = TRUE)
```



Gradient Boosted Trees

APPLICATIONS



OF DATA SCIENCE

Boosting: The Gist

Compute the average response \bar{y} , and use this as the initial predicted value for each sample

For $j = 1$ to m do

- Compute the residual for each observation
- Sample a fraction of the original data
- Fit a regression tree of a certain "*interaction depth*" using the residual as response
- Predict each sample using this tree
- Update the predicted value of each sample by adding the previous iteration's value to current predicted value X *learning rate* or *shrinkage* parameter

End.

What you end up with:

$$\hat{y}_i = \bar{y} + \sum_{j=1}^m \lambda \cdot tree_j(\mathbf{x}_i)$$

That is, GBT (or GBM) in general is called an additive model.

💡 Where does "Gradient" come from?

- The added random sampling of the data was added later, this version is called *Stochastic* Gradient Boosting, the default for the parameter often called "bagging fraction" is 0.5
- The shrinkage λ can be turned into λ_j , a unique weight for each added tree
- Note the **many** tuning parameters here. Do you know how to tune multiple params?
- Not so parallel now, eh?

GBT with gbm

```
library(gbm)


mod_gbm_na <- gbm(income ~ ., data = okupid2_train,
                     distribution = "gaussian",
                     n.trees = 500, shrinkage = 0.1, bag.fraction = (
pred_gbm_na <- predict(mod_gbm_na, okupid2_valid, n.trees = 500)
report_rmse_and_cor(okupid2_valid$income, pred_gbm_na)
```

```
## RMSE: 0.351
## CORR: 0.504
```

```
mod_gbm_mice <- gbm(income ~ ., data = okupid2_imp_mice_train,
                      distribution = "gaussian", n.trees = 500, shrinkage =
pred_gbm_mice <- predict(mod_gbm_na, okupid2_imp_mice_valid, n.trees = 500)
report_rmse_and_cor(okupid2_valid$income, pred_gbm_mice)
```

```
## RMSE: 0.353
## CORR: 0.497
```

For variables importance see the `summary` function:



Do you see the difference in the profile of importance between RF and GBT?

GBT with xgboost

xgboost is faster and has some additional features, e.g. the ability to train with a validation set until it shows no improvement, similar to neural networks.

```
library(xgboost)


dtrain <- xgb.DMatrix(data = data.matrix(okcupid2_train[, predictors],
                                         label = okcupid2_train$income)
dval <- xgb.DMatrix(data = data.matrix(okcupid2_valid[, predictors],
                                         label=okcupid2_valid$income))
watchlist <- list(train=dtrain, test=dval)

mod_xgboost_na <- xgb.train(data = dtrain, watchlist=watchlist,
                               objective = "reg:squarederror", nrounds = 1000,
                               eta = 0.1, subsample = 0.5,
                               early_stopping_rounds = 10, verbose = TRUE)

pred_xgboost_na <- predict(mod_xgboost_na, dval)
report_rmse_and_cor(okcupid2_valid$income, pred_xgboost_na)
```

```
## RMSE: 0.358
## CORR: 0.474
```

For variables importance see the `xgb.importance` function:



The CART (Classification)

APPLICATIONS



OF DATA SCIENCE

Detour: A Classification Problem

APPLICATIONS



OF DATA SCIENCE

OKCupid: Predicting Dogs or Cats People

It won't be easy:

```
okcupid %>% count(pets)

## # A tibble: 16 x 2
##   pets                  n
##   <chr>                <int>
## 1 dislikes cats          122
## 2 dislikes dogs           44
## 3 dislikes dogs and dislikes cats 196
## 4 dislikes dogs and has cats     81
## 5 dislikes dogs and likes cats   240
## 6 has cats                1406
## 7 has dogs                 4134
## 8 has dogs and dislikes cats   552
## 9 has dogs and has cats      1474
## 10 has dogs and likes cats    2333
## 11 likes cats               1063
## 12 likes dogs                7224
## 13 likes dogs and dislikes cats 2029
## 14 likes dogs and has cats     4313
## 15 likes dogs and likes cats   14814
## 16 <NA>                   19921
```

We will define cats/dogs people in the following way and stick to non-NA observations:

```
cats_categories <- c("has cats", "likes cats", "dislikes dogs and  
                      "dislikes dogs and has cats")  
dogs_categories <- c("has dogs", "likes dogs", "has dogs and dislikes  
                      "likes dogs and dislikes cats")  
okcupid3 <- okcupid %>%  
  mutate(pets = case_when(  
    pets %in% cats_categories ~ "cats",  
    pets %in% dogs_categories ~ "dogs",  
    TRUE ~ NA_character_)) %>%  
  drop_na(pets)  
  
okcupid3 %>% count(pets)
```

```
## # A tibble: 2 x 2  
##   pets      n  
##   <chr> <int>  
## 1 cats     2550  
## 2 dogs    13939
```

Notice the classes are very unbalanced, with roughly 11 dogs people to every 2 cats people.

As before in the vector predictors we have 36 continuous and categorical variables which may or may not be predictive to being a cats/dogs person:

```
okupid3 <- okupid3 %>%
  select(pets, predictors) %>%
  mutate(id = 1:n())

dim(okupid3)
```

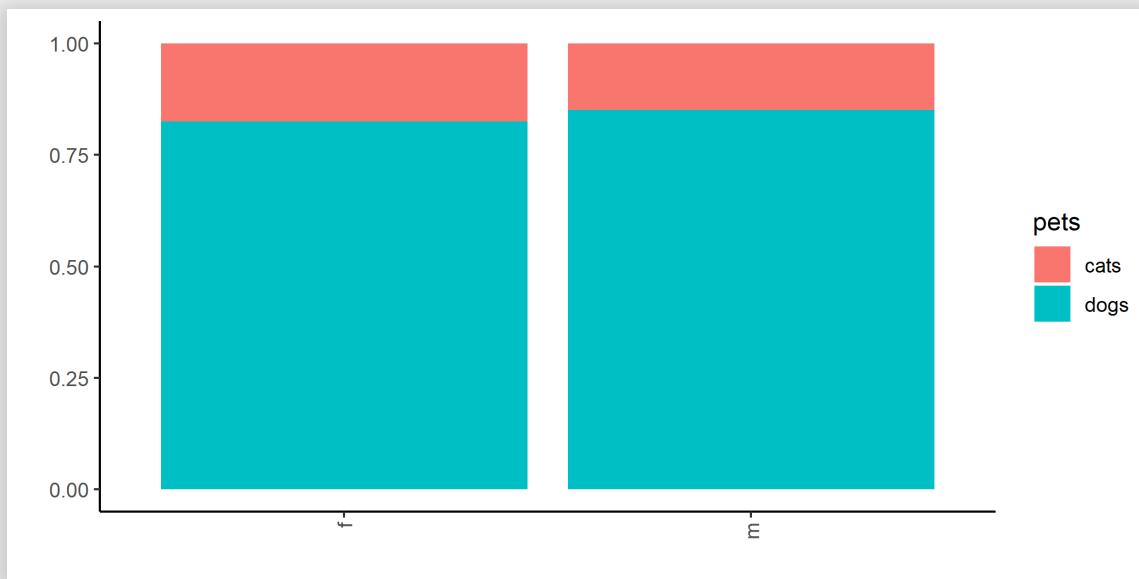
```
## [1] 16489     38
```

And as before we split the data into training, test and validation sets, not shown here:

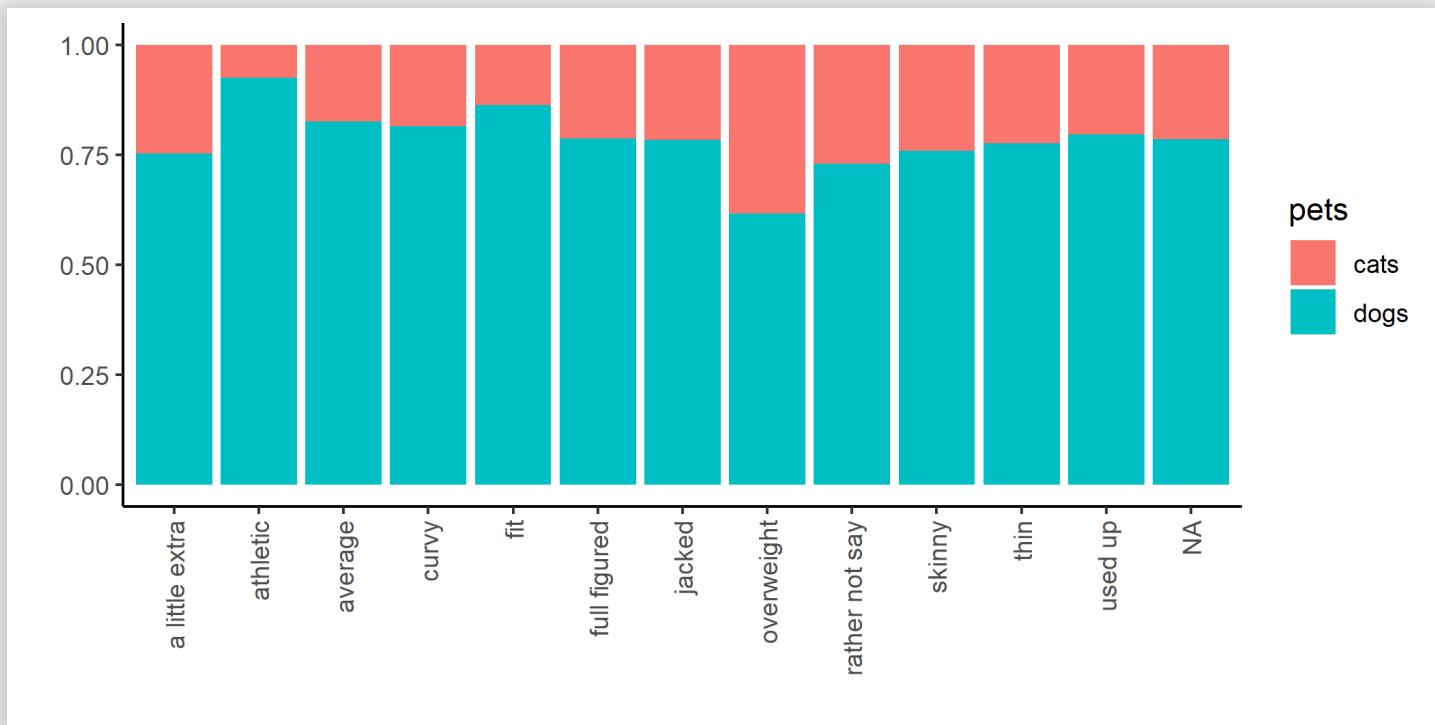
```
## train no. of rows: 10000
## validation no. of rows: 2489
## test no. of rows: 4000
```

Worth exploring relations between predictors and being a cats/dogs person:


```
var_vs_pets_stackedbar <- function(var) {  
  ggplot(okcupid3_train, aes({{var}}, fill = pets)) +  
  geom_bar(position = "fill") +  
  theme_classic() +  
  labs(x = "", y = "") +  
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust =  
})  
  
var_vs_pets_stackedbar(sex)
```




```
var_vs_pets_stackedbar(body_type)
```



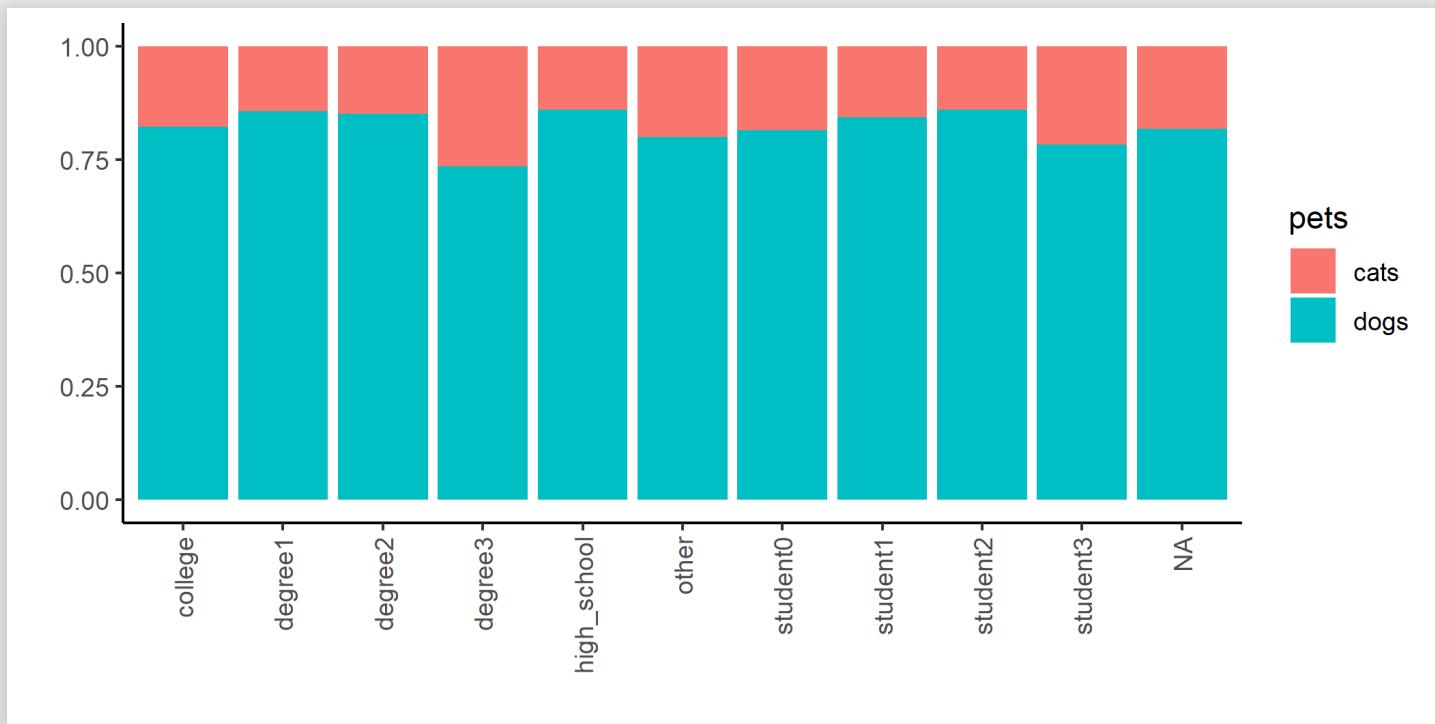
```
var_vs_pets_stackedbar(sign2)
```



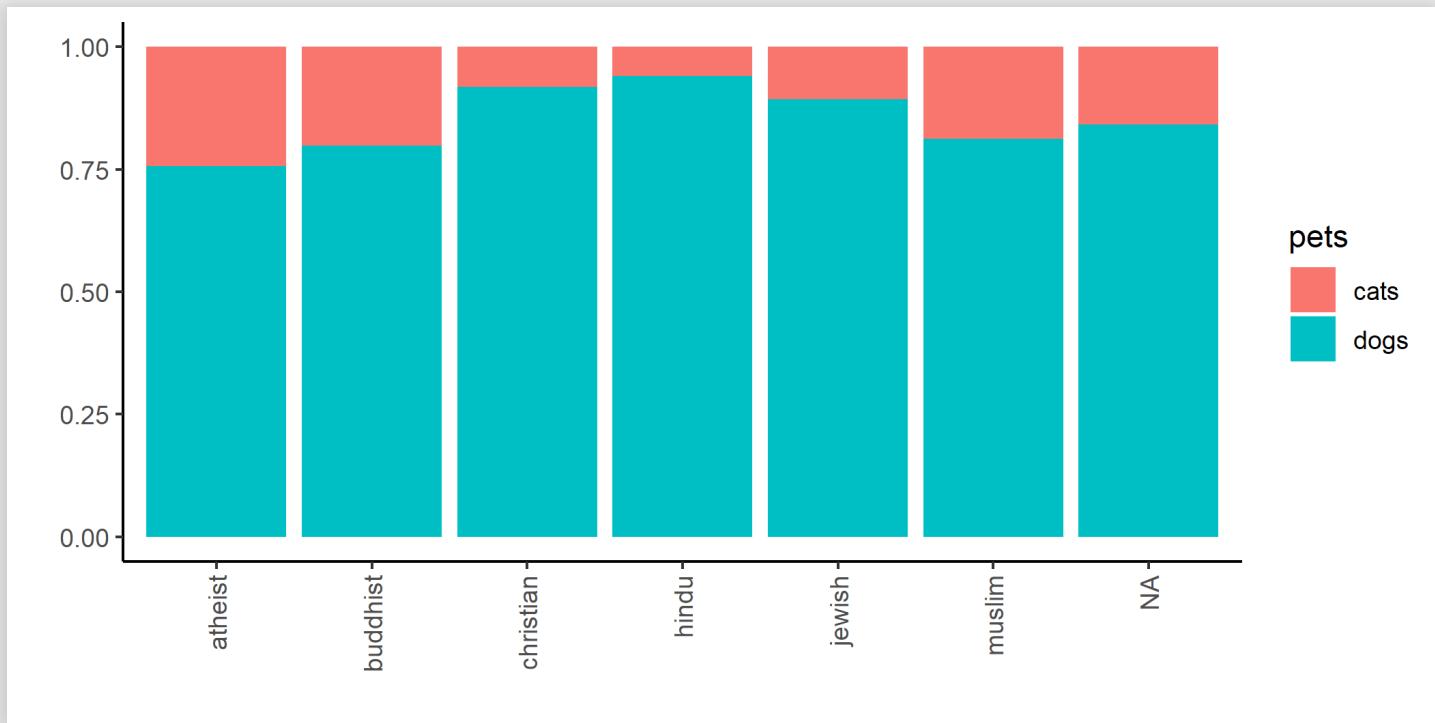
```
var_vs_pets_stackedbar(job3)
```



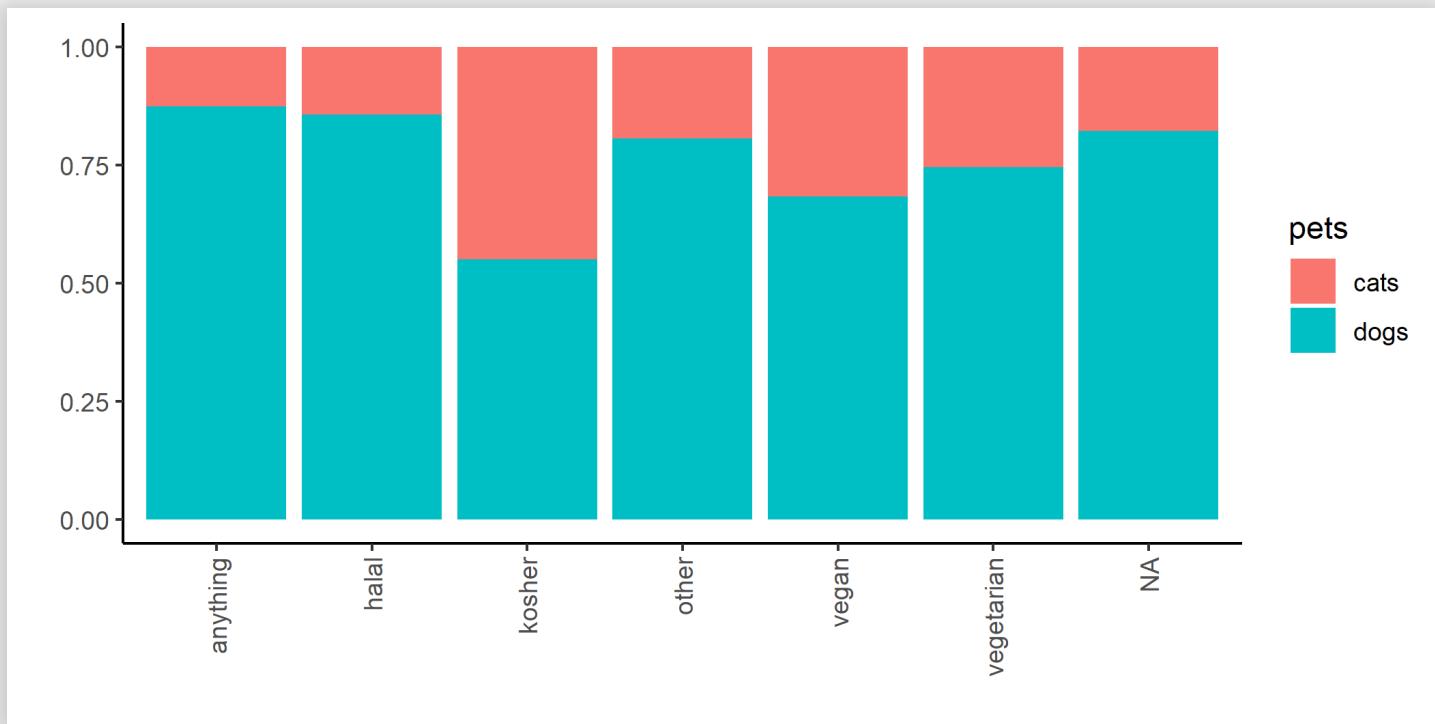
```
var_vs_pets_stackedbar(education2)
```



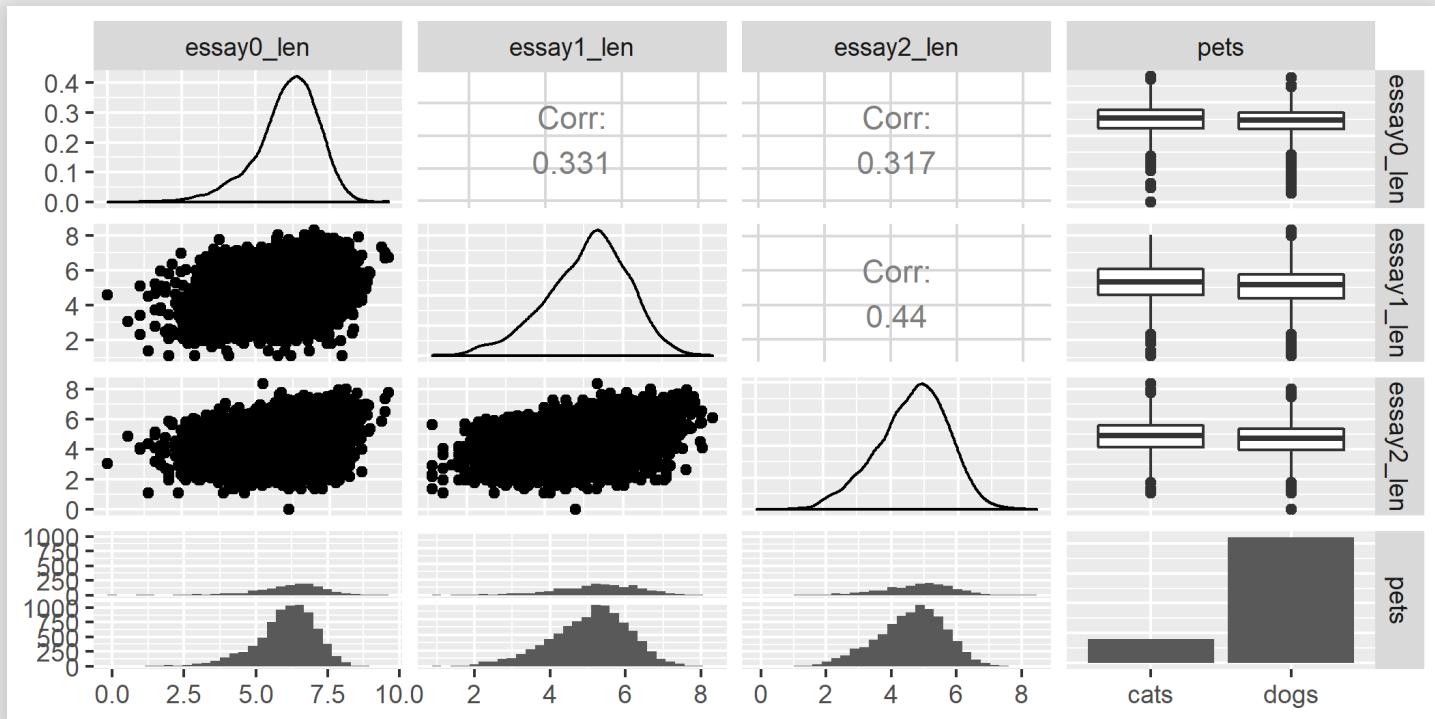
```
var_vs_pets_stackedbar(religion2)
```



```
var_vs_pets_stackedbar(diet2)
```



```
ggpairs(okcupid3_train %>%
         select(essay0_len:essay2_len, pets))
```



Baseline: Logistic Regression

As in `lm`, it is better to impute missing values than to lose these observations:

```
# library(mice)
# mice_obj <- mice(okcupid3, m = 1, maxit = 10, seed = 42)
# okcupid3_imp_mice <- complete(mice_obj)

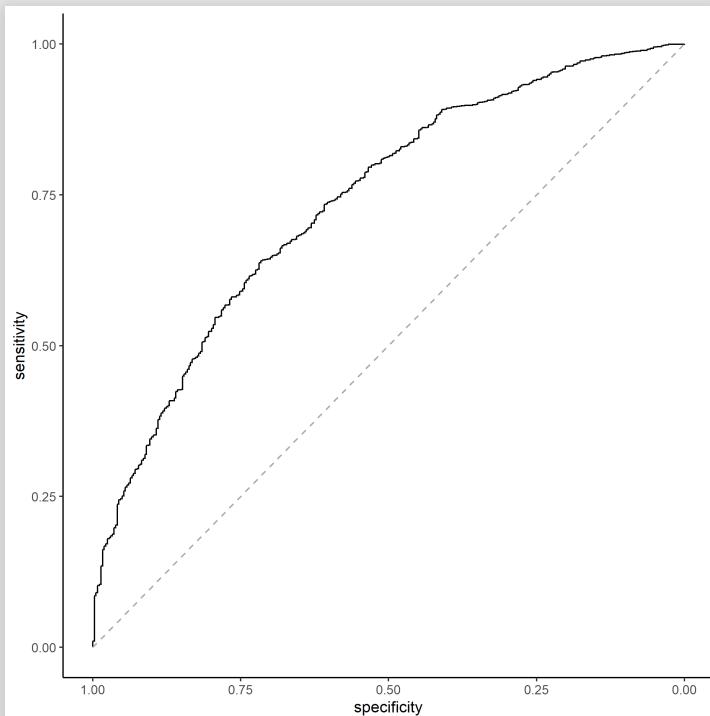
okcupid3_imp_mice <- read_rds("../data/okcupid3_imp_mice.rds")
okcupid3_imp_mice_train <- okcupid3_imp_mice[train_idx, ]
okcupid3_imp_mice_valid <- okcupid3_imp_mice[valid_idx, ]

mod_glm_mice <- glm(pets ~ ., data = okcupid3_imp_mice_train, family = "binomial")
pred_glm_mice <- predict(mod_glm_mice, okcupid3_imp_mice_valid, type = "response")
```

Let's plot the ROC curve:

```
library(pROC)
roc_obj <- roc(okcupid3_valid$pets, pred_glm_mice)

ggroc(roc_obj) +
  theme_classic() +
  geom_segment(aes(x = 1, xend = 0, y = 0, yend = 1), color="darkgrey")
```



Choosing a specific cutoff, the proportion of dogs people in the training sample (the "positive" class), 0.84, let's look at the confusion matrix:

```
cutoff <- mean(okcupid3_train$pets == "dogs")
pred_class <- ifelse(pred_glm_mice > cutoff, "dogs", "cats")
true_class <- okcupid3_valid$pets

table(true_class, pred_class)
```

```
##           pred_class
## true_class  cats  dogs
##       cats    238   125
##       dogs    678 1448
```

💡 What would be the accuracy? Recall (sensitivity, specificity)? Precision (PPV, NPV)?

We would like to know the model's AUC, accuracy for a given score cutoff, recall and precision:

```
report_accuracy_and_auc <- function(obs, pred, cutoff = mean(okcupy$pet_type == "Dogs")) {  
  roc_obj <- roc(obs, pred)  
  AUC <- as.numeric(auc(roc_obj))  
  res <- coords(roc_obj, x = cutoff,  
                ret = c("accuracy", "recall", "precision", "specificity"),  
                transpose = TRUE)  
  glue("AUC: {format(AUC, digits = 3)}  
ACC: {format(res['accuracy'], digits = 3)}  
Dogs: Recall: {format(res['recall'], digits = 3)}  
      Precision: {format(res['precision'], digits = 3)}  
Cats: Recall: {format(res['specificity'], digits = 3)}  
      Precision: {format(res['npv'], digits = 3)}")  
}  
  
report_accuracy_and_auc(okcupid3_valid$pets, pred_glm_mice)
```

```
## AUC: 0.738  
## ACC: 0.677  
## Dogs: Recall: 0.681  
##           Precision: 0.921  
## Cats: Recall: 0.656  
##           Precision: 0.26
```

Baseline: GLMNET: Penalized Logistic Regression

```
okcupid3_imp_mat_train <- model.matrix(~ ., okcupid3_imp_mice_train)
okcupid3_imp_mat_valid <- model.matrix(~ ., okcupid3_imp_mice_validation)

glmnet_cv <- cv.glmnet(x = okcupid3_imp_mat_train,
                        y = okcupid3_train$pets, family = "binomial")

best_lambda <- glmnet_cv$lambda.min

mod_glm_glmnet <- glmnet(x = okcupid3_imp_mat_train,
                           y = okcupid3_train$pets,
                           family = "binomial", lambda = best_lambda)
pred_glm_glmnet <- predict(mod_glm_glmnet, okcupid3_imp_mat_valid,
                           type = "response")

report_accuracy_and_auc(okcupid3_valid$pets, pred_glm_glmnet[, 1])
```

```
## AUC: 0.739
## ACC: 0.672
## Dogs: Recall: 0.672
##           Precision: 0.923
## Cats: Recall: 0.672
##           Precision: 0.259
```

End of Detour

APPLICATIONS



OF DATA SCIENCE

The OG CART

1. Find the predictor to (binary) split on and the value of the split, by *Gini* (a.k.a impurity) criterion
2. For each resulting node if $n_{node} > 20$ go to 1
3. Once full tree has been grown, perform pruning using the *cost-complexity parameter* c_p and the $Gini_{c_p}$ criterion
4. Predict the proportion of each class at each terminal node, or most common class

The *Gini* criterion

- y is the discrete dependent variable with J classes
- The $Gini_{parent}$ at a parent node is:
$$\sum_{j=1}^J P(y = j)[1 - P(y = j)] = \sum_j p_j(1 - p_j)$$
- a continuous predictor v is nominated for splitting the current node with splitting value l
- such that S_1 is the set of observations for which $v_i \leq l$
- and S_2 is the set of observations for which $v_i > l$
- p_{1j} is the (observed) probability of y equals class j in set S_1

$$Gini_{split} = \frac{\#\text{obs node1}}{\#\text{obs parent}} Gini_{node1} + \frac{\#\text{obs node2}}{\#\text{obs parent}} Gini_{node2} = \\ \frac{\#\text{obs node1}}{\#\text{obs parent}} \sum_j p_{1j}(1 - p_{1j}) + \frac{\#\text{obs node2}}{\#\text{obs parent}} \sum_j p_{2j}(1 - p_{2j})$$

And we find l for which $\Delta G = Gini_{parent} - Gini_{split}$ is greatest.


Why Gini?

For $J = 2$ classes this means:

$$p_1(1 - p_1) + p_2(1 - p_2) = 2p_1(1 - p_1)$$


The more "impure" the node, the higher this metric:

```
p <- seq(0, 1, 0.01)
plot(p, 2 * p * (1 - p), type = "l")
```



For example if religion2 is candidate in splitting pets:

```
gini <- function(l, df, v) {  
  pets_above <- df %>% filter({{v}} >= l) %>% pull(pets)  
  pets_below <- df %>% filter({{v}} < l) %>% pull(pets)  
  p_dogs_above <- mean(pets_above == "dogs")  
  gini_above <- 2 * p_dogs_above * (1 - p_dogs_above)  
  p_dogs_below <- mean(pets_below == "dogs")  
  gini_below <- 2 * p_dogs_below * (1 - p_dogs_below)  
  gini_weighted <- (length(pets_above) / nrow(df)) * gini_above +  
    (length(pets_below) / nrow(df)) * gini_below  
  return(gini_weighted)  
}  
  
age <- seq(18, 69, 1)  
gini_age <- map_dbl(age, gini, df = okcupid3_train, v = age)  
  
p1 <- okcupid3_train %>%  
  group_by(age) %>%  
  summarise(p_dogs = mean(pets == "dogs"), n = n()) %>%  
  ggplot(aes(age, p_dogs)) +  
  geom_point(aes(size = n)) +  
  theme_classic() +  
  labs(x = "") +  
  theme(axis.text.x = element_blank())  
  
p2 <- tibble(age = age, gini = gini_age) %>%  
  ggplot(aes(age, gini)) +  
  geom_line() +
```



CART with rpart

Let's tune c_p for our data using a 5-fold (manual) Cross Validation.
The criterion to maximize would be AUC.

```
n_cv <- 5; cp_seq <- seq(0, 0.01, 0.0005)


okcupid3_train_val <- okcupid3_train %>%
  mutate(val = sample(1:n_cv, n(), replace = TRUE))

get_validation_set_auc <- function(i, .cp) {
  ok_tr <- okcupid3_train_val %>% filter(val != i) %>% select(-val)
  ok_val <- okcupid3_train_val %>% filter(val == i) %>% select(-val)
  mod <- rpart(pets ~ ., data = ok_tr,
                control = rpart.control(cp = 0, xval = 1))
  mod <- prune(mod, cp = .cp)
  pred <- predict(mod, ok_val) [, 2]
  roc_obj <- roc(ok_val$pets, pred)
  as.numeric(auc(roc_obj))
}

get_cv_auc <- function(.cp) {
  tibble(cp = rep(.cp, n_cv),
        auc = map_dbl(1:n_cv, get_validation_set_auc, .cp = .cp))
}
```

```
cv_table <- map_dfr(cp_seq, get_cv_auc)

cv_table %>%
  mutate(cp = factor(cp)) %>%
  ggplot(aes(cp, auc)) +
  geom_boxplot() +
  theme_classic() +
  labs(x = "") +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust =
```




Training on the entire training set:

```
mod_tree_na_class <- rpart(pets ~ ., data = okupid3_train,  
                           control = rpart.control(cp = 0, xval =  
mod_tree_na_class <- prune(mod_tree_na_class, cp = 0.0015)
```

There's no way to plot this tree nicely...

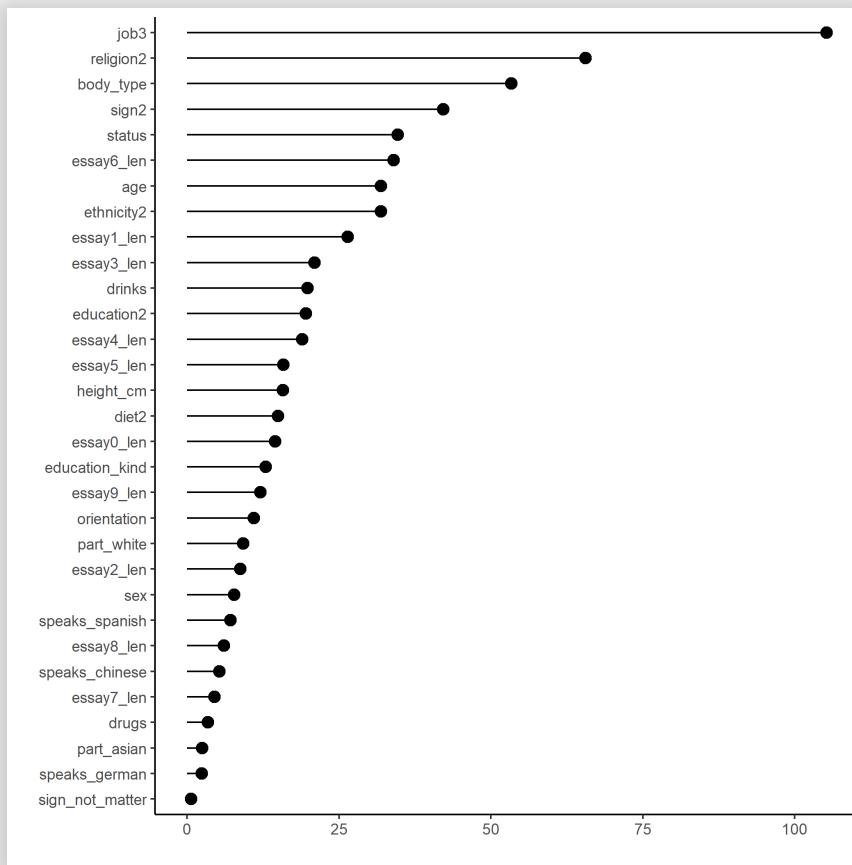
```
plot(mod_tree_na_class)
```



Printing the tree it seems like after splitting observations by religion, those who are Christian, Jewish or Hindu have a 88% of being dogs people, and the rest are then split by Status, Body Type, the length of their 9th Essay, etc...

```
Fitted party:  
[1] root  
| [2] religion2 in atheist, buddhist, muslim  
| [3] status in available, married, seeing someone  
| | [4] body_type in a little extra, curvy, full figured, skinny, thin  
| | | [5] job3 in administrative, craftsmanship, hardware, health, management, musical, other, tech  
| | | | [6] essay3_len >= 4.45447: cats (n = 35, err = 11.4%)  
| | | | [7] essay3_len < 4.45447  
| | | | | [8] sign2 in aquarius, cancer, capricorn, virgo: cats (n = 9, err = 11.1%)  
| | | | | [9] sign2 in aries, gemini, leo, pisces, sagittarius, scorpio: dogs (n = 15, err = 26.7%)  
| | | | [10] job3 in academia, financial, government, legal services, marketing, media, rather not say, student, transportation, travel: dogs  
[11] body_type in athletic, average, fit, jacked, overweight, rather not say, used up  
| [12] diet2 in kosher, vegan, vegetarian  
| | [13] essay0_len < 6.77815: cats (n = 30, err = 23.3%)  
| | | [14] essay0_len >= 6.77815: dogs (n = 16, err = 31.2%)  
| [15] diet2 in anything, halal, other  
| | [16] essay8_len >= 5.30087  
| | | [17] essay2_len < 5.59658: cats (n = 17, err = 17.6%)  
| | | | [18] essay2_len >= 5.59658: dogs (n = 27, err = 37.0%)  
| | [19] essay8_len < 5.30087  
| | | [20] essay4_len >= 7.88766: cats (n = 9, err = 22.2%)  
| | | | [21] essay4_len < 7.88766: dogs (n = 150, err = 24.7%)  
[22] status in single  
| [23] body_type in a little extra, average, curvy, fit, full figured, jacked, overweight, skinny, thin, used up  
| | [24] ethnicity2 in native american, other, white  
| | | [25] drinks in desperately, not at all, rarely, very often  
| | | | [26] job3 in academia, administrative, hardware, legal services, management, media, other, rather not say, tech, transportation  
| | | | | [27] education2 in degree3, other, student0, student1  
| | | | | | [28] body_type in a little extra, average, full figured, jacked, overweight, skinny, thin  
| | | | | | | [29] job3 in academia, administrative, hardware, media, rather not say, tech: cats (n = 39, err = 12.8%)  
| | | | | | | [30] job3 in legal services, other: dogs (n = 7, err = 28.6%)  
| | | | | [31] body_type in curvy, fit, used up: dogs (n = 27, err = 33.3%)  
| | | | [32] education2 in college, degree1, degree2, high_school, student1, student2  
| | | | [33] body_type in a little extra, curvy, full figured: cats (n = 16, err = 25.0%)  
| | | | [34] body_type in average, fit, jacked, overweight, skinny, thin, used up  
| | | | | [35] education2 in degree1, degree2, student2  
| | | | | | [36] sign2 in capricorn, gemini, libra, pisces, taurus  
| | | | | | | [37] essay5_len < 5.02025  
| | | | | | | | [38] job3 in hardware, legal services, management, media: cats (n = 16, err = 18.8%)  
| | | | | | | | [39] job3 in academia, other, tech: dogs (n = 13, err = 38.5%)  
| | | | | | | | [40] essay5_len >= 5.02025: dogs (n = 8, err = 25.0%)  
| | | | | [41] sign2 in aquarius, aries, cancer, leo, sagittarius, scorpio, virgo
```

Variables Importance



Prediction

```
pred_tree_na_class <- predict(mod_tree_na_class, okcupid3_valid)

report_accuracy_and_auc(okcupid3_valid$pets, pred_tree_na_class[,]

## AUC: 0.612
## ACC: 0.788
## Dogs: Recall: 0.87
##           Precision: 0.881
## Cats: Recall: 0.309
##           Precision: 0.289
```

As expected, far from impressive.

Let's try using the imputed NA data:

```
mod_tree_imp_class <- rpart(pets ~ ., data = okcupid3_imp_mice_train,
                             control = rpart.control(cp = 0, xval = 10))
mod_tree_imp_class <- prune(mod_tree_imp_class, cp = 0.0015)

pred_tree_imp_class <- predict(mod_tree_imp_class, okcupid3_imp_mice_train)

report_accuracy_and_auc(okcupid3_valid$pets, pred_tree_imp_class[, 1])

## AUC: 0.672
## ACC: 0.748
## Dogs: Recall: 0.793
##           Precision: 0.9
## Cats: Recall: 0.485
##           Precision: 0.285
```